

**CONTENTS :**

1. Aggregation Function
2. Basic SQL Function
3. Collect Statistics
4. Data Manipulation Language (DML)
5. Date Functions
6. Distinct Vs Group By Functions
7. Explain
8. Format Functions
9. Help and Show
10. Join Functions
11. Join Indexes
12. Math Functions
13. OLAP Functions
14. Substrings and Positioning Functions
15. Temporal Tables Create function
16. Temporary Tables
17. Teradata Parallel Transport
18. The Quantile Function
19. Top SQL Command Cheat Sheet
20. View Functions
21. The Where Clause
22. Sample
23. Set Operators functions
24. Statistical Aggregate Functions
25. Stored Procedure Functions
26. Sub Query Functions



**Set 1 :**

- 1) Aggregation Function
- 2) Basic SQL Function
- 3) Collect Statistics
- 4) Data Manipulation Language (DML)
- 5) Date Functions



## Aggregation Function



### Aggregation Function

#### Overview

"Teradata climbed Aggregate Mountain and delivered a better way to Sum It."  
- Tera-Tom Coffing

Quiz – You calculate the Answer Set in your own Mind

Aggregation_Table	
Employee_No	Salary
423400	100000.00
423401	100000.00
423402	NULL

```
SELECT AVG(Salary) as "AVG"  
      ,Count(Salary) as SalCtn  
      ,Count(*)       as RowCtn  
FROM Aggregation_Table ;
```

AVG	SalCtn	RowCtn
-----	--------	--------

Please fill in the  
values you  
think will be  
in the Answer.

What would the result set be from the above query? The next slide shows answers!

Answer – You calculate the Answer Set in your own Mind

Aggregation_Table	
Employee_No	Salary
423400	100000.00
423401	100000.00
423402	NULL

```
SELECT AVG(Salary) as "AVG"  
      ,Count(Salary) as SalCtn  
      ,Count(*)       as RowCtn  
FROM Aggregation_Table ;
```



AVG	SalCnt	RowCnt
100000.00	2	3

Here are  
all the  
Correct  
answers

Here are your answers!

#### The 3 Rules of Aggregation

Aggregation_Table	SELECT AVG(Salary), Count(Salary), Count(*) FROM Aggregation_Table;
Employee_No      Salary	
423400	100000.00
423401	100000.00
423402	NULL

- 1) Aggregates Ignore **Null** Values.
- 2) Aggregates WANT to come back in **one** row.
- 3) You CAN'T mix Aggregates with **normal columns** unless you use a **GROUP BY**.

AVG(Salary) = \$100000.00	Count(Salary) = 2	Count(*) = 3
------------------------------	----------------------	-----------------

There are Five Aggregates



**Aggregation\_Table**

Employee_No	Salary
423400	100000.00
423401	100000.00
423402	NULL

There are **FIVE AGGREGATES** which are the following:

**MIN** – The Minimum Value.

**MAX** – The Maximum Value.

**AVG** – The Average of the Column Values.

**SUM** – The Sum Total of the Column Values.

**COUNT** – The Count of the Column Values.

**Quiz – How many rows come back?**

Employee_No	Dept_No	Last_Name	First_Name	Salary
12345678	100	Cheney	Mande	48950.00
1235249	400	Harrison	Peter	54500.00
2341218	400	Reilly	William	36000.00
2312225	300	Larkins	Lorraine	40200.00
2000000	?	Jones	Squiggy	32800.50
1000034	10	Smythe	Richard	36000.00
1123454	40	Stockling	James	54500.00
1324657	200	Coffing	Billy	41881.88
1333454	200	Smith	John	48000.00



Query →

```
SELECT MIN(Salary)
      ,MAX(Salary)
      ,SUM(Salary)
      ,AVG(Salary)
      ,Count(*)  
FROM Employee_Table;
```

How many rows will the above query produce in the result set? The answer is one.

**Troubleshooting Aggregates**

Employee_Table				
Employee_No	Dept_No	Last_Name	First_Name	Salary
123456789	100	Chandira	Michelle	54500.00
123563459	400	Tarrison	Sherrett	54500.00
23412118	400	Reilly	William	36000.00
2312225	300	Larkins	Lorraine	40200.00
2000000	?	Jones	Squiggy	32800.00
1000000	10	White	Ronald	54500.00
1121334	400	Stickling	Cletus	54500.00
1324657	200	Coffing	Billy	41898.88
1333454	200	Smith	John	48000.00



```
SELECT Dept_No → [NON-Aggregate]
      .MIN(Salary)
      .MAX(Salary)
      .SUM(Salary)
      .AVG(Salary)
      .COUNT(*)
FROM Employee_Table;
```

How many rows will the above query produce in the result set? None, because this query errors!

GROUP BY when Aggregates and Normal Columns Mix

Employee_No	Dept_No	Last_Name	First_Name	Salary
12345678	100	Chambers	Mandee	45850.00
12345679	400	Brown	Heath	14500.00
12341218	400	Reilly	William	36000.00
12312225	300	Larkins	Lorraine	40200.00
12345670	300	James	Susan	32000.00
12000234	10	Smythe	Richard	22800.00
11113334	400	Stickling	Cletus	54500.00
12344547	200	Coffing	Billy	41880.00
12334545	200	Smith	John	48000.00

```
SELECT Dept_No ← [NON-Aggregate]
      .MIN(Salary)
      .MAX(Salary)
      .SUM(Salary)
      .AVG(Salary)
      .COUNT(*)
FROM Employee_Table
GROUP BY Dept_No;
```

A GROUP BY statement is needed when mixing aggregates and non-aggregates.

GROUP BY Delivers one row per Group



NON-Aggregate

```

SELECT Dept_No
      ,MIN(Salary)
      ,MAX(Salary)
      ,SUM(Salary)
      ,AVG(Salary)
      ,Count(*)
FROM Employee_Table
GROUP BY Dept_No
ORDER by 1;
  
```

The Group BY Dept\_No allows for the Aggregates to be calculated per Dept\_No.

GROUP BY Dept\_No or GROUP BY 1 the same thing

<pre> SELECT Dept_No       ,MIN(Salary)       ,MAX(Salary)       ,SUM(Salary)       ,AVG(Salary)       ,Count(*) FROM Employee_Table GROUP BY Dept_No ORDER BY Dept_No;   </pre>
--

<pre> SELECT Dept_No       ,MIN(Salary)       ,MAX(Salary)       ,SUM(Salary)       ,AVG(Salary)       ,Count(*) FROM Employee_Table GROUP BY 1 ORDER BY 1;   </pre>
--

Both queries above produce the same result. The GROUP BY allows you to either name the column or use the number in



the SELECT list, just like the ORDER BY.

#### Limiting Rows and Improving Performance with WHERE

Employee Table				
Employee_No	Dept_No	Last_Name	First_Name	Salary
1323578	100	Chambers	Mandeep	48500.00
1256349	400	Harrison	Herbert	54500.00
2341218	400	Reilly	William	36000.00
2311225	300	Jarkins	Siraj	38000.00
1200000	10	Jones	Suganya	32800.50
1000234	10	Smythe	Richard	32800.00
1121334	400	Stickling	Cletus	54500.00
1324657	200	Coffing	Billy	41988.88
1333454	200	Smith	John	48000.00

```
SELECT Dept_No, MIN(Salary), MAX(Salary), SUM(Salary)
   , AVG(Salary), COUNT(*)
FROM Employee Table
WHERE Dept_No IN (200, 400) ← WHERE Clause acts
GROUP BY Dept_No           as a filter before any
Order by 1;                  Calculations are done
```

Will Dept\_No 300 be calculated? Of course you know it will... NOT!

WHERE Clause in Aggregation limits unneeded Calculations

```
SELECT Dept_No, MIN(Salary), MAX(Salary), SUM(Salary)
   , AVG(Salary), COUNT(*)
FROM Employee Table
WHERE Dept_No IN (200, 400) ← WHERE Clause acts
GROUP BY Dept_No           as a filter before any
Order by 1;                  Calculations are done
```

The system eliminates reading any other Dept\_No other than 200 and 400. This means that only the Dept\_No of 200 and 400 will come off the disk to be calculated.

Keyword HAVING tests Aggregates after they are Totaled



```
SELECT Dept_No, MIN(Salary), MAX(Salary), SUM(Salary)
     ,AVG(Salary), COUNT(*)
  FROM Employee_Table WHERE Dept_No in (200, 400)
 GROUP BY Dept_No
 HAVING Count(*) > 2 ;
```

HAVING Clause acts as a filter on all Aggregates after they are totaled.

Previous Answer Set				
Dept_No	Min(Salary)	Max(Salary)	Sum(Salary)	Avg(Salary)
200	41888.00	48000.00	89888.00	44944.44
400	38000.00	54500.00	148000.00	48333.33

NEW Answer Set  
???????????????

Can you calculate what the new Answer Set will be after the HAVING Clause is implemented?

The HAVING Clause only works on Aggregate Totals. The WHERE filters rows to be excluded from calculation, but the HAVING filters the Aggregate totals after the calculations, thus eliminating certain Aggregate totals.

Keyword HAVING is like an Extra WHERE Clause for Totals

```
SELECT Dept_No, MIN(Salary), MAX(Salary), SUM(Salary)
     ,AVG(Salary), COUNT(*)
  FROM Employee_Table WHERE Dept_No in (200, 400)
 GROUP BY Dept_No
 HAVING Count(*) > 2 ;
```

Having Clause acts as a filter on all Aggregates after they are totaled.

Previous Answer Set without the HAVING Statement				
Dept_No	Min(Salary)	Max(Salary)	Sum(Salary)	Avg(Salary)
200	41888.00	48000.00	89888.00	44944.44
400	38000.00	54500.00	145000.00	48333.33



## Teradata CoE

## Teradata Lab Course

New Answer Set using the HAVING Statement

Dept_No	MinSalary	MaxSalary	SumSalary	AvgSalary	Count(1)
400	36000.00	54900.00	145800.00	48333.33	3

The HAVING Clause only works on Aggregate Totals after they are totaled. It is a final check after aggregation is complete.  
Now only the totals with Count() > 2 can return.

Getting the Average Values per Column

```
SELECT 'Product_ID' AS "Column Name"  
,COUNT(DISTINCT(product_id)) AS "Average Rows"  
FROM Sales_Table ;
```

Column Name	Average Rows
Product_ID	3

```
SELECT 'Sale_Date' AS "Column Name"  
,COUNT(*) / COUNT(DISTINCT(Sale_Date)) AS "Average Rows"  
FROM Sales_Table ;
```

Column Name	Average Rows
Sale Date	3

The first query retrieved the average rows per value for the column Product\_ID. The example below did the same for the column Sale\_Date.

Average Values per Column for All Columns in a Table

```
SELECT 'Product_ID' AS "Column Name"  
,COUNT(DISTINCT(product_id)) AS "Average Rows"  
'Sale Date' AS "Column Name"  
,COUNT(*) / COUNT(DISTINCT(Sale_Date)) AS "Average Rows"  
FROM Sales_Table ;
```

Column Name	Average Rows	Column Name	Average Rows
Product ID	3	Sale Date	3

The query above retrieved the average rows per value for both columns in the table.

Three types of Advanced Grouping



Sales_Table		
Product_ID	Sale_Date	Daily_Sales
1000	2000-08-28	41888.88
2000	2000-08-28	41888.88
3000	2000-08-28	41881.77
1000	2000-08-29	41888.88
2000	2000-08-29	48000.00
3000	2000-08-29	48000.00
1000	2000-08-30	36000.07
2000	2000-08-30	49349.84
3000	2000-08-30	49349.84
1000	2000-09-01	40205.43
2000	2000-09-01	40205.43
3000	2000-09-01	28000.00
1000	2000-09-02	36021.93
2000	2000-09-02	36021.93
3000	2000-09-02	18478.94

Be prepared to be amazed. There are three advanced options listed above for grouping data. Each is more powerful than the one before. The following pages will give great examples.

GROUP BY Grouping Sets

```
SELECT Product_ID
      ,EXTRACT (MONTH FROM Sale_Date) AS MTH
      ,EXTRACT (YEAR FROM Sale_Date) AS YR
      ,SUM(Daily_Sales) AS SUM_Daily_Sales
   FROM Sales_Table
 GROUP BY GROUPING SETS (Product_ID, MTH, YR)
 ORDER BY Product_ID Desc, MTH Desc, YR Desc;
```

Product_ID	MTH	YR	SUM_Daily_Sales
1000	?	?	223487.88
2000	?	?	506014.31
1000	?	?	331294.72
?	10	?	445634.99
?	09	?	418769.36
?	?	2000	862404.35

GROUP BY GROUPING Sets above will show you what your Daily\_Sales were for each Product\_ID, for each month, and for each year.

GROUP BY Rollup



```

SELECT Product_ID
      ,EXTRACT(MONTH FROM Sale_Date) AS MTH
      ,EXTRACT(YEAR FROM Sale_Date) AS YR
      ,SUM(Daily_Sales) AS SUM_Daily_Sales
  FROM Sales_Table
 GROUP BY ROLLUP (Product_ID, MTH, YR)
 ORDER BY Product_ID Desc, MTH Desc, YR Desc;
  
```

Product_ID	MTH	YR	SUM_Daily_Sales
3000	10	2000	84908.06
3000	10	?	84908.06
3000	9	2000	139779.76
3000	9	?	139779.76
3000	?	?	224587.82
2000	10	2000	166872.90
2000	10	?	166872.90
2000	9	2000	139738.91
2000	9	?	139738.91
2000	?	?	306611.81
1000	10	2000	191854.03
1000	10	?	191854.03
1000	9	2000	139350.69
1000	9	?	139350.69
1000	?	?	331204.72
?	?	?	862404.35

**GROUP BY Rollup Result Set**

Product_ID	MTH	YR	SUM_Daily_Sales
3000	10	2000	84908.06
3000	10	?	84908.06
3000	9	2000	139779.76
3000	9	?	139779.76
3000	?	?	224587.82
2000	10	2000	166872.90
2000	10	?	166872.90
2000	9	2000	139738.91
2000	9	?	139738.91
2000	?	?	306611.81
1000	10	2000	191854.03
1000	10	?	191854.03
1000	9	2000	139350.69
1000	9	?	139350.69
1000	?	?	331204.72
?	?	?	862404.35

This is the full result set from the previous GROUP BY ROLLUP query.

**GROUP BY Cube**

```

SELECT Product_ID
      ,EXTRACT(MONTH FROM Sale_Date) AS MTH
      ,EXTRACT(YEAR FROM Sale_Date) AS YR
      ,SUM(Daily_Sales) AS SUM_Daily_Sales
  FROM Sales
 GROUP BY CUBE(Product_ID, MTH, YR)
 ORDER BY Product_ID Desc, MTH Desc, YR Desc;
  
```

Product_ID	MTH	YR	SUM_Daily_Sales
3000	10	2000	84908.06
3000	9	2000	136479.76
3000	?	?	189799.00
3000	?	2000	224587.52
3000	?	?	224587.52
2000	10	2000	168972.90
2000	10	?	168972.90

GROUP BY ROLLUP displays Daily\_Sales for each Product\_ID, for each distinct month, for each month per year, for each year, plus a grand total.

GROUP BY CUBE Result Set



**Teradata CoE**      **Teradata Lab Course**

Product_ID	MTH	YR	SUM_Daily_Sales	
3000	10	2000	84908.06	Prod 1000 October, 2000 sales
3000	10	?	84908.06	Prod 3000 October sales all years
3000	9	2000	139679.76	Prod 3000 September, 2000 sales
3000	9	?	139679.76	Prod 3000 September sales all years
3000	?	2000	224587.82	Prod 3000 sales for the year 2000
3000	?	?	224587.82	Prod 3000 sales for all years
2000	10	2000	166872.90	Prod 2000 October, 2000 sales
2000	10	?	166872.90	Prod 2000 October sales all years
2000	9	2000	139728.91	Prod 2000 September, 2000 sales
2000	?	2000	306611.81	Prod 2000 sales for the year 2000
2000	?	?	306611.81	Prod 2000 sales for all years
1000	10	2000	191854.03	Prod 1000 October, 2000 sales
1000	10	?	191854.03	Prod 1000 October sales all years
1000	9	2000	139350.69	Prod 1000 September, 2000 sales
1000	9	?	139350.69	Prod 1000 September sales all years
1000	?	2000	331204.72	Prod 1000 sales for the year 2000
1000	?	?	331204.72	Prod 1000 sales for all years
?	10	2000	443634.99	Total Sales for October, 2000
?	10	?	443634.99	Total Sales for all October dates
?	9	2000	418769.36	Total Sales for September, 2000
?	9	?	418769.36	Total Sales for all September dates
?	?	2000	862404.35	Total Sales for the year 2000
?	?	?	862404.35	Total Sales for all years totalled

Use the Nexus for all Groupings



In Nexus, just right click on the Sales\_ Table and choose Super Join Builder. Then, select all the columns. Then, choose the Analytics tab on the top right. Choose Grouping Sets in the Analytics Tab. Then, drag the Product\_ID column to the Product. Drag the Sale\_Date to the Date Column. Then, drag the Daily\_Sales column to the Sum. Then, Check Box all the Group BY Functions on the right of the screen and then hit Execute or Send SQL to Nexus, and you are Done!



## Testing Your Knowledge – Basic Aggregation

Employee_Table				
Employee_No	Dept_No	Last_Name	First_Name	Salary
1232578	100	Chen	Wendy	49850.00
1256349	400	Harrison	Herbert	54500.00
2341218	400	Reilly	William	36000.00
2312225	300	Larkins	Lorraine	40200.00
2000000	?	Jones	Squiggy	32800.50
1000004	400	Smythe	Ronald	32500.00
1121334	400	Stockling	Cleatus	54500.00
1324657	200	Coffing	Billy	41888.88

First, SELECT the AVERAGE Salary from the Employee\_Table.

## Testing Your Knowledge – Multiple Aggregates

Employee_Table				
Employee_No	Dept_No	Last_Name	First_Name	Salary
1232578	100	Chen	Wendy	49850.00
1256349	400	Harrison	Herbert	54500.00
2341218	400	Reilly	William	36000.00
2312225	300	Larkins	Lorraine	40200.00
2000000	?	Jones	Squiggy	32800.50
1000004	400	Smythe	Ronald	32500.00
1121334	400	Stockling	Cleatus	54500.00
1324657	200	Coffing	Billy	41888.88
1324654	200	Smith	John	48000.00

Now, SELECT the AVERAGE Salary and the SUM of the Salary from the Employee\_Table.

## Testing Your Knowledge - Group By

Employee_Table				
Employee_No	Dept_No	Last_Name	First_Name	Salary
1232578	100	Chen	Wendy	49850.00
1256349	400	Harrison	Herbert	54500.00
2341218	400	Reilly	William	36000.00
2312225	300	Larkins	Lorraine	40200.00
2000000	?	Jones	Squiggy	32800.50
1000004	400	Smythe	Ronald	32500.00
1121334	400	Stockling	Cleatus	54500.00
1324657	200	Coffing	Billy	41888.88
1324654	200	Smith	John	48000.00

Now, SELECT the AVERAGE Salary and the SUM of the Salary from the Employee\_Table PER DEPARTMENT (Dept\_No).



## Testing Your Knowledge - Using a Where Clause

Employee_No	Dept_No	Last_Name	First_Name	Salary
1123279	100	Chase	Maudie	48500.00
1256349	400	Harrison	Herbert	54500.00
2341214	800	Reilly	William	36000.00
2312225	300	Larkins	Lorraine	40200.00
2000000	?	Jones	Squiggy	32800.50
1000234	10	Smith	Richard	32800.00
1121334	400	Stockling	Cleatus	54500.00
1324657	200	Coffing	Billy	41888.88
1333454	200	Smith	John	48000.00

After that, SELECT the AVERAGE Salary and the SUM of the Salary from the Employee\_Table PER DEPARTMENT (Dept\_No). However, I only want to see the people from Department 200, 300, 400.

## Testing Your Knowledge-Using Having

Employee_No	Dept_No	Last_Name	First_Name	Salary
1123279	100	Chase	Maudie	48500.00
1256349	400	Harrison	Herbert	54500.00
2341218	400	Reilly	William	36000.00
2312225	300	Larkins	Lorraine	40200.00
2000000	?	Jones	Squiggy	32800.50
1000234	10	Smith	Richard	32800.00

1121334	400	Stockling	Cleatus	54500.00
1324657	200	Coffing	Billy	41888.88
1333454	200	Smith	John	48000.00

After that, SELECT the AVERAGE Salary and the SUM of the Salary from the Employee\_Table PER DEPARTMENT (Dept\_No). However, I only want to see Department 200, 300, 400 which has an AVERAGE Salary of over 43,000.

## Final Answer to Test Your Knowledge on Aggregates

Answer Set		
Dept_No	AvgSalary	SumSalary
200	44944.44	39988.00
400	43314.67	144500.00

```
Select Dept_No, AVG(Salary), SUM(Salary)
From Employee_Table
Where Dept_No IN (200, 300, 400)
Group By Dept_No
Having AVG(Salary) > 43000
```



This should be your final answer set. The query under it should be approximately what you wrote to attain such an answer set. How'd you do?

## Basic SQL Functions



### Basic SQL Functions

#### Introduction

Student_Table				
Student_ID	Last_Name	First_Name	Class_Code	Grade_Pt
423400	Larkins	Michael	FR	0.00
123456	Baker	John	FR	2.00
280023	McRoberts	Richard	JR	1.50
240000	Johnson	Stanley	S	7
1234567	White	Tom	SO	3.80
124111	Thomas	Wendy	FR	4.50
124222	Davis	David	SR	3.25
123280	Phillips	Martin	SR	3.00
321139	Bout	Jimmy	JR	2.45
333456	Smith	Andy	SO	2.00

The Student\_Table above will be used

In our examples.

This is a portion of the Student\_Table which we will use to present some basic examples of SQL and get some hands-on

experience with querying this table. This block attempts to show you the table, show you the query, and show you the

result set.



## Teradata CoE

## Teradata Lab Course

### SELECT \* (All Columns) in a Table

The screenshot shows the Teradata Studio interface with a query window containing the following SQL code:

```
SELECT *  
FROM Students;
```

The results pane displays the following data from the Students table:

Student_ID	First_Name	Last_Name	Class_Code	Grade_Pt
1_A	423400	Adele	Michael	8.00
2_A	281222	Wilson	Steve	3.80
3_A	280300	John	Robert	3.90
4_A	321233	Dawn	James	3.85
5_A	128414	Roseann	Kerry	3.80
6_A	323300	Andy	Lee	2.00
7_A	324612	Delivery	Darren	3.35
8_A	261000	Julianne	Wendy	1
9_A	234240	Wendy	TR	4.00
10_A	121250	Millette	Alma	3.00

Most every SQL statement will consist of a SELECT and a FROM. You SELECT the columns you want to see on your report, and an Asterisk (\*) means you want to see all columns in the table on the returning answer set!

### SELECT Specific Columns in a Table



```

1 SELECT First_Name
2   ,Last_Name
3   ,Class_Code
4   ,Grade_Pt
5 FROM Student_Table;

```

Column names must be separated by commas. Notice that only the columns requested come back on the report, not all columns. Also, notice that the order of the columns in the SQL is the same order on the report.

Using the Best Form for Writing SQL.

Student_Table					
Student_ID	Last_Name	First_Name	Class_Code	Grade_Pt	
423400	Larkins	Michael	FR	0.00	
125634	Hanson	Henry	FR	2.88	
280023	McRoberts	Richard	JR	1.90	
260000	Johnson	Stanley	?	?	
234121	Wilson	Susie	SO	3.00	
234121	Tracy	Wendy	FR	4.00	
324652	Delaney	Danny	SR	3.35	
132350	Phillips	Martin	SR	3.00	
324652	Brown	Jeffrey	SR	3.50	
333450	Smith	Andy	SO	3.00	

SELECT First_Name,	SELECT First_Name	Can you
Last_Name,	Class_Code	see the
Class_Code,	,Class_Code	difference
Grade_Pt	,Grade_Pt	between
FROM Student_Table;	FROM Student_Table;	these two examples?

Why is the example on the right better even though they are functionally equivalent?



## Commas in the Front or in the Back?

1      `SELECT First_Name, Last_Name, Class_Code, Grade_Pt  
FROM Student_Table;]`

2      `SELECT First_Name, Last_Name, Class_Code, Grade_Pt  
FROM Student_Table;]`

Michael    Hawkins    PR    9.00  
Henry      Hanson     PR    2.00  
Richard    McRoberts    JR    1.00  
Sue        Tracy        SO    7.00  
Susie      O'Wilson    SO    3.00  
Darryl     Thomas      PR    4.00  
Danny      Delaney    SR    3.50  
Martin     Phillips    SR    5.00  
Randy      Rogers     SR    1.00  
Andy       Smith       SO    2.00

Commas in the front (example 1) is Tera-Tom's recommendation to writing, but the next page is an even better example for a company standard. Both queries will produce the same answer set and have the same performance.

Place your Commas in front for better Debugging Capabilities

`SELECT First_Name,  
Last_Name,  
Class_Code,  
Grade_Pt  
FROM Student_Table;  
Error!`

Sometimes if  
you Add or  
Delete a  
COLUMN you  
can overlook an  
ending Comma!

`SELECT First_Name  
Last_Name  
Class_Code  
Grade_Pt  
FROM Student_Table;`

Successful

Having commas in front to separate column names makes it easier to debug.

## Sort the Data with the ORDER BY Keyword



**Teradata CoE**      **Teradata Lab Course**

**Netezza Query Commander**

File Edit View Insert Tools Window Help  
 Home Database Teradata SQL CLASSIC Help  
 New Query = New Query -> ETL -> Dashboard -> Indicator -> History

SQL DDL DML Reporting Utilities System Vertica

**SELECT \*  
FROM Student\_Table  
ORDER BY Last\_Name;**

Student ID	Last Name	First Name	Class	Code	Grade Pt
322133	Bond	Jimmy	JR		3.95
324652	Delaney	Danny	SR		3.35
125634	Hausner	Henry	FR		2.88
260000	Johnson	Stanley	?	?	
423400	Larkins	Michael	FR		0.00
280023	McRoberts	Richard	JR		1.90
123250	Phillips	Martin	SR		3.00
333459	Smith	Andy	SO		2.00
234121	Thomas	Wendy	FR		4.00
231222	Wilson	Sue	SO		3.80

Rows typically come back to the report in random order. To order the result set, you must use an ORDER BY. When you order by a column, it will order in ASCENDING order. This is called the Major Sort!

**ORDER BY Defaults to Ascending**

Sorts the Answer Set In Ascending Order By Last\_Name

**SELECT \*  
FROM Student\_Table  
ORDER BY Last\_Name;**

Student ID	Last Name	First Name	Class	Code	Grade Pt
322133	Bond	Jimmy	JR		3.95
324652	Delaney	Danny	SR		3.35
125634	Hausner	Henry	FR		2.88
260000	Johnson	Stanley	?	?	
423400	Larkins	Michael	FR		0.00
280023	McRoberts	Richard	JR		1.90
123250	Phillips	Martin	SR		3.00
333459	Smith	Andy	SO		2.00
234121	Thomas	Wendy	FR		4.00
231222	Wilson	Sue	SO		3.80

When you use the ORDER BY statement, it will default to ascending order. But you can change that if you like. I will show you how to do that in a few pages.

Use the Name or the Number in your ORDER BY Statement

---

 Capgemini  
CONSULTING. INNOVATION. OUTSOURCING

**Teradata CoE**      **Teradata Lab Course**

The ORDER BY can use a number to represent the sort column. The number 2 represents the second column on the report. This is also going to default to ascending.

The ORDER BY can use a number to represent the sort column. The number 2 represents the second column on the report. This is also going to default to ascending.

**Two Examples of ORDER BY using Different Techniques**

SELECT * FROM Student_Table ORDER BY 1;	SELECT * FROM Student_Table ORDER BY Grade_Pt;
Student_ID Last_Name First_Name Class_Code Grade_Pt 1 221331 Bond Bonny JR 3.95 2 324652 Delaney Danny SR 3.35 3 123434 Hanson Henry FR 2.88 4 260934 Phillips Andy FR 3.00 5 421400 Larkins Michael FR 0.00 6 280921 McRoberts Richard JR 1.90 7 123456 Phillips Linda FR 3.00 8 333450 Smith Andy SO 2.00 9 234121 Thomas Wendy FR 4.00 1 211222 Wilson Jessie SO 3.80	Student_ID Last_Name First_Name Class_Code Grade_Pt 1 221331 Bond Bonny JR 3.95 2 324652 Delaney Danny SR 3.35 3 123434 Hanson Henry FR 2.88 4 260934 Phillips Andy FR 3.00 5 421400 Larkins Michael FR 0.00 6 280921 McRoberts Richard JR 1.90 7 123456 Phillips Linda FR 3.00 8 333450 Smith Andy SO 2.00 9 234121 Thomas Wendy FR 4.00 1 211222 Wilson Jessie SO 3.80

You have the option of using a number instead of the column name. The columns number is represented by what position it is in the SELECT statement, not the table. If you use an \* in your Select Statement, then the columns number is represented by the position it is in the table. The two above queries are the same.

**Changing the ORDER BY to Descending Order**

---

**Teradata CoE**

**Teradata Lab Course**

The screenshot shows the Nexus Query Channeler interface. At the top, there's a toolbar with File, Edit, View, Tools, Web, Windows, Help, and various icons for Execute, New Query, ETL, Dashboard, Visualize, and History. Below the toolbar, the title bar says "Teradata CoE" and "Teradata Lab Course". The main area has a tree view on the left labeled "Systems" with nodes like Oracle, DB2, MySQL, PostgreSQL, MySQL, PostgreSQL, and Vertica. A central pane titled "Query 1" contains the following SQL code:

```

SELECT *
FROM Student_Table
ORDER BY Last_Name DESC;
    
```

Below the code, a message box says "Results 1" and "Drop a column header here to group by that column." The results table shows the following data:

	Student_ID	Last_Name	First_Name	Class_Code	Grade_Pt
1	322133	Bond	Jimmy	JR	3.95
2	324652	Delaney	Danny	SR	3.35
3	125634	Hansen	Henry	FR	2.88
4	260900	Johnson	Stanley	?	?
5	423400	Larkins	Michael	FR	0.00
6	280023	McRoberts	Richard	JR	1.90
7	123250	Phillips	Martin	SR	3.00
8	333450	Smith	Andy	SO	2.00
9	234121	Thomas	Wendy	FR	4.00
1	231222	Wilson	Susie	SO	3.80

At the bottom of the interface, there's a Capgemini logo.

## Teradata CoE

## Teradata Lab Course

If you want the data sorted in descending order just place DESC at the end:

NULL Values sort First in Ascending Mode (Default)

```
SELECT *           SELECT *
FROM Student_Table FROM Student_Table
ORDER BY 5;        ORDER BY Grade_Pt;
```

Student ID	Last Name	First Name	Class Code	Grade Pt
250000	Johansson	Stanley	?	?
423400	Lakins	Michael	FR	0.00
280023	McRoberts	Richard	JR	1.90
333450	Hanson	Andy	SO	2.00
123534	Hanson	Henry	FR	2.88
123250	Phillips	Mariu	SR	3.00
324652	Delaney	Danny	SR	3.35
231222	Wilson	Susie	SO	3.80
322133	Bond	Jimmy	JR	3.95
234121	Thomas	Wendy	FR	4.00

Nulls

The default for ORDER BY is Ascending mode (ASC). Notice that this places the Null Values at the beginning of the Answer Set.

NULL Values sort Last in Descending Mode (DESC)

```
SELECT *           SELECT *
FROM Student_Table FROM Student_Table
ORDER BY 5 DESC;   ORDER BY Grade_Pt DESC;
```

Student ID	Last Name	First Name	Class Code	Grade Pt
234121	Thomas	Wendy	FR	4.00
322133	Bond	Jimmy	JR	3.95
231222	Wilson	Susie	SO	3.80
324652	Delaney	Danny	SR	3.35
123250	Phillips	Marin	SR	3.00
123450	Hanson	Henry	FR	2.88
333450	Hanson	Andy	SO	2.00
280023	McRoberts	Richard	JR	1.90
423400	Lakins	Michael	FR	0.00
250000	Johansson	Stanley	?	?

Nulls are

Last in

Order

You can ORDER BY in descending order by putting a DESC after the column name or its corresponding number. Null Values will sort Last in DESC order.



## Teradata CoE

## Teradata Lab Course

### Major Sort vs. Minor Sorts

```
SELECT * FROM Student_Table  
ORDER BY Class_Code DESC,  
        Grade_Pt ASC;
```

Major Sort on  
Class\_Code DESCENDING

Minor Sort on  
Grade\_Pt Ascending

Student_ID	Last Name	First Name	Class Code	Grade Pt
12320	Phillips	Marta	SR	Major
32145	DeLoach	Danny	SR	Secs
331450			SO	Firs
231222	Wilson	Steve	SO	In
280023	McRoberts	Richard	JR	DESC
322133	Bond	January	JR	Order
321421	Anderson	Patricia	FR	4.95
125634	Hanson	Honey	FR	2.88
234121	Thomas	Wendy	FR	4.00
280009	Johnson	Stanley	?	?

Major sort is how things are sorted, but a minor sort kicks in if there are Major Sort ties.

### Multiple Sort Keys using Names vs. Numbers

```
SELECT Employee_No  
      ,Last_Name  
      ,First_Name  
      ,Dept_No  
      ,Salary  
  FROM Employees_Table  
 ORDER BY Last_Name  
      ,First_Name  
      ,Dept_No  
      ,Salary DESC;
```

These queries sort identically

Queries can have a multiple columns in the ORDER BY. The first column in an ORDER BY is called the MAJOR SORT.

Those after it are MINOR SORTS.

Both of these Queries do the same thing. Once they sort Dept\_No column in DESC order, they'll sort any ties by LAST\_NAME. If any ties still occur, they'll sort by SALARY. Let me show you a real world example in the next slide!

### Sorts are Alphabetical, NOT Logical

```
SELECT * FROM Student_Table  
ORDER BY Class_Code;
```



Student_ID	Last_Name	First_Name	Class_Code	Grade_Pt
246000	Johnson	Stanley	F	3
234121	Thomas	Wendy	FR	4.00
125634	Hanson	Henry	FR	2.88
423400	Larkins	Michael	FR	0.00
322133	Bond	Jimmy	JR	3.95
322134	Murder	Richard	JR	1.50
331222	Wilcox	Susie	SO	3.80
333450	Smith	Andy	SO	2.00
324652	Delaney	Danny	SR	3.35
123250	Phillips	Martin	SR	3.00

This sorts alphabetically, but Sophomores (SO) logically come after Freshman

Change the query to Order BY Class\_Code logically (FR, SO, JR, SR, ?)

Using A CASE Statement to Sort Logically

```

SELECT * FROM Student_Table
ORDER BY CASE Class_Code
          WHEN 'FR' THEN 1
          WHEN 'SO' THEN 2
CASE in the
          WHEN 'JR' THEN 3
ORDER BY
          WHEN 'SR' THEN 4
Statement
          ELSE 5
END;

```



**Teradata CoE****Teradata Lab Course**

Student_ID	Last_Name	First_Name	Class_Code	Grade_Pt
234121	Thomas	Wendy	FR	4.00
125614	Hanson	Henry	FR	2.88
423400	Larkins	Michael	FR	0.00
333450	Smith	Andy	SO	2.00
231222	Wilson	Susie	SO	3.80
322133	Bondo	Jimmy	JR	3.95
280923	McRoberts	Richard	JR	1.90
324652	Delaney	Danny	SR	3.35
123250	Phillips	Martin	SR	3.00
260000	Johnson	Stanley	?	?

We are using a CASE Statement to Order BY Class\_Code logically (FR, SO, JR, SR.)

How to ALIAS a Column Name

```
SELECT First_Name AS Fname
      ,Last_Name  AS Lname
      ,Class_Code AS "Class Code"
      ,Grade_Pt   AS "AVG"
      ,Student_ID AS STU_ID
  FROM Student_Table;
```

ALIAS Rules!

- 1) AS is optional
- 2) Use Double Quotes when Spaces are in the Alias
- 3) Use Double Quotes when the Alias is a reserved word

Different Techniques  
for Aliasing

When you ALIAS a column, you give it a new name for the report header. You should always reference the column using the ALIAS everywhere else in the query. You never need Double Quotes in SQL unless you are Aliasing.

A Missing Comma can by Mistake become an Alias

```
SELECT First_Name,Last_Name,Class_Code,Grade_Pt
  FROM Student_Table;
```

Missing a Comma



First Name	Last Name	Grade_Pt
Michael	Larkins	FR
Sue	Wilson	SO
Richard	McRoberts	JR
Jimmy	Bond	JR
Henry	Harrison	FR
Ashley	Santos	SO
Danny	Delaney	SR
Stanley	Johnson	?
Wendy	Thomas	FR
Martin	Phillips	SR

Column names must be separated by commas. Notice in this example, there is a comma missing between Class\_Code and Grade\_Pt. This will result in only three columns appearing on your report, and one of them will be titled wrong.

#### The Title Command and Literal Data

```
SELECT 'Character Data'
      , 'Character Data' (TITLE 'Character//Data')
      , 123 (TITLE 'Numeric Data')
      , 'Character Data' (TITLE 'My//Stacked//Example');
```

Stacks the Report Header

Character Data	Character Data	Numeric Data	MY Stacked Example
Character Data	Character Data	123	Character Data

A Literal Value brings back the Literal Value! Also, notice that the word 'Character' is stacked over the 'Data' portion of the heading for the second column using the Nexus Query Character. So, as an alternative, a TITLE can be used instead of an alias, if you want to stack the words in the output title.

The difference between an ALIAS and a TITLE is that the ALIAS can be used in the SQL again, such as in the ORDER BY or WHERE statements. But, a TITLE is only good for the report heading. Notice that Title uses Single Quotes not double quotes.

Comments using Double Dashes are Single Line Comments

Comment → -- Double Dashes provide a single line comment
 

```
SELECT *
  FROM Student_Table
  ORDER BY Grade_Pt;
```



**Teradata CoE****Teradata Lab Course**

Student_ID	Last_Name	First_Name	Class_Code	Grade_Pt
260000	Johnson	Stanley	S	3
423400	Larkins	Michael	FR	0.00
280023	McRoberts	Richard	JR	1.90
333450	Smith	Andy	SO	2.00
125634	Hanson	Henry	FR	2.88
123250	Phillips	Maria	SR	3.00
324652	Delaney	Danny	SR	3.35
231222	Wilson	Susie	SO	3.80
322133	Bond	Jimmy	JR	3.95
234121	Thomas	Wendy	FR	4.00

Double dashes make a single line comment that will be ignored by the system.

**Comments for Multi-Lines**

Comment → /\* This is how you can make multi-line comments to express what is going on in the code. \*/

```
SELECT *  
FROM Student_Table  
ORDER BY Grade_Pt;
```

Student_ID	Last_Name	First_Name	Class_Code	Grade_Pt
260000	Johnson	Stanley	S	3
423400	Larkins	Michael	FR	0.00
280023	McRoberts	Richard	JR	1.90
333450	Smith	Andy	SO	2.00
125634	Hanson	Henry	FR	2.88
123250	Phillips	Maria	SR	3.00
324652	Delaney	Danny	SR	3.35
231222	Wilson	Susie	SO	3.80
322133	Bond	Jimmy	JR	3.95
234121	Thomas	Wendy	FR	4.00

/\\* starts a multi-line comment, and \*\



## Teradata CoE                                  Teradata Lab Course

### Comments for Multi-Lines as Double Dashes per Line

**Comments** → -- This is how you can make multi-line comments  
 -- also to express what is going on in the code.

```
SELECT *  
FROM Student_Table  
ORDER BY Grade_Pt;
```

Student_ID	Last_Name	First_Name	Class_Code	Grade_Pt
260000	Johnson	Stanley	3	3
423400	Larkins	Michael	FR	0.00
280023	McRoberts	Richard	JR	1.90
319111	Smith	Andy	S0	2.00
125634	Benson	Henry	FR	2.85
123250	Phillips	Martin	SR	3.00
324652	Delaney	Danny	SR	3.35
231122	Wilson	Sue	S0	3.45
322133	Bond	Jimmy	SR	3.95
234121	Thomas	Wendy	FR	4.00

You can make multi-line comments with double dashes on each line.

### A Great Technique for Comments to Look for SQL Errors

**SELECT** Student\_ID  
 .Last\_Name  
 .First\_Name  
 .Class\_Code as Sum  
 .Grade\_Pt  
**FROM** Student\_Table  
**WHERE** Grade\_Pt > 3.6

**Comment** → -- SELECT Student\_ID  
 .Last\_Name  
 .First\_Name  
 .Class\_Code as Sum  
 .Grade\_Pt  
**FROM** Student\_Table  
**WHERE** Grade\_Pt > 3.6

Student_ID	Last_Name	First_Name	Grade_Pt
234121	Thomas	Wendy	4.00
231122	Wilson	Sue	3.80
322133	Bond	Jimmy	3.95

Sometimes you get an error in your SQL, and it is difficult to find. When our first query ran, it produced an error. We were not sure if our Class\_Code was the error, so we commented that line out and ran our SQL again. Everything ran perfectly the next time, so we knew the Class\_Code line must have been the



error. What was the error? The alias Sum is a reserved word. Comments can be used to test lines for errors.



## Collect Statistics

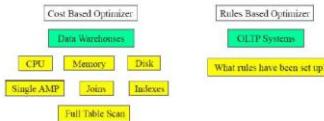


## Collect Statistics

### The Teradata Parsing Engine (Optimizer) is Cost Based

The Parsing Engine (PE) is often referred to as the Teradata Optimizer, and it will actually generate several plans to choose from and ultimately choose the one with the lowest cost of resources. This is critical to performance in supporting mixed workloads ranging from OLTP to large joins and Decision Support (DS). All cost based optimizers require statistical information about the data and the machine resources (CPU, disk, memory).

The other type of optimizer is a rules based optimizer which is designed for transactional On-Line Transaction Processing (OLTP) workloads where queries are well known and the data has been logically and physically structured to support OLTP workloads.



A cost based optimizer is much better than an optimizer that is rule based for data warehouses.

### The Purpose of Collect Statistics

The Teradata Parsing Engine (PE) is in charge of creating the PLAN for the AMPs to follow. The PE works best when Statistics have been collected on a table. Then it knows:

1. The number of rows in the table
2. The average row size
3. Information on all Indexes in which statistics were collected
4. The range of values for the column(s) in which statistics were collected
5. The number of rows per value for the column(s) in which statistics were collected
6. The number of NULLs for the column(s) in which statistics were collected

The purpose of the COLLECT STATISTICS command is to gather and store demographic data for one or more columns or indices of a table or join index. This process computes a statistical profile of the collected data and stores the synopsis in the Data Dictionary (DD) inside USER\_DBC for use during the PE's optimizing phase of SQL statement parsing. The optimizer uses this synopsis data to generate efficient table access and join plans. Do NOT COLLECT Statistics on all columns in the table.

When Teradata Collects Statistics, it creates a Histogram



Histogram of Employee Table Column Last_Name			
Interval 1	Interval 2	Interval 3	Interval 4
1 Axled	Bubaker	Custer	Dzabo
2 Anderson	Bell	Caner	Davis
3 55	150	50	100
4 100	200	300	50
5 900	800	900	800

When statistics are collected, Teradata does a full table scan, sorts the column or index, places them into a default of 250 intervals, and then provides the above Histogram.

1. Highest Sort Value in the Interval
2. Most Frequent Value in the Interval
3. Rows with the Most Frequent Value
4. Other Values in the Interval
5. Number of Rows of other Values

This is what is stored in statistics. This is tricky to understand at first, but recognize first that there are 55 people with a Last\_Name of Anderson, 150 Bells, 50 Canes, and 100 with the name Davis. Each interval shows the most popular value and row count.

#### The Interval of the Collect Statistics Histogram

Histogram		Anyone who has a Last Name that is between 'A' and 'Axled' falls into this interval
1	Axled	1
2	Anderson	2
3	55	3
4	100	4
5	900	5

There are 900 rows other than Anderson in this interval.

The PE now knows there are 55 Andersons in the table, and it assumes, for any other name falling between 'A' and 'Anderson', that there are 9 values for each ( $900 / 100 = 9$ )

When statistics are collected, Teradata does a full table scan, sorts the column or index, places them into a default of 250 intervals, and then provides the above Histogram. This is what the PE uses to build a plan. Above, you see only interval one of 250.

#### Histogram Quiz



Histogram of Employee_Table Column Last_Name				
Interval 1	Interval 2	Interval 3	Interval 4	
1 Axelrod	Burbaker	Custer	Dzoba	
2 Anderson	<b>Bell</b>	Cane	Davis	
3 55	150	50	160	
4 100	200	300	50	
5 900	800	900	800	

1. Highest Sorted Value in the Interval
  2. Most Frequent Value in the Interval
  3. Rows with the Most Frequent Value
  4. Other Values in the Interval
  5. Number of Rows of other Values
1. Which Interval would the PE look to find the Last\_Name of 'Apple'?\_\_\_\_\_
  2. How many people are in the Employee\_Table with a Last\_Name of 'Davis'?\_\_\_\_\_
  3. In Interval 2 how many other names are there other than 'Bell'?\_\_\_\_\_
  4. How many people named 'Baker' would Teradata estimate?\_\_\_\_\_
  5. How many people name 'Donaldson' would Teradata estimate?\_\_\_\_\_
  6. How many people named 'Cooper' would Teradata estimate?\_\_\_\_\_

Answers to Histogram Quiz



Histogram of Employee_Table Column Last_Name				
	Interval 1	Interval 2	Interval 3	Interval 4
1	Axelrod	Burbaker	Custer	Dzoba
2	Anderson	Bell	Cane	Davis
3	55	150	50	160
4	100	210	300	50
5	900	800	900	800

1. Highest Sorted Value in the Interval
  2. Most Frequent Value in the Interval
  3. Rows with the Most Frequent Value
  4. Other Values in the Interval
  5. Number of Rows of other Values
1. Which Interval would the PE look to find the Last\_Name of 'Apple'? **Interval 1**
  2. How many people are in the Employee\_Table with a Last\_Name of 'Davis'? **160**
  3. In Interval 2 how many other names are there other than 'Bell'? **200 other names**
  4. How many people named 'Baker' would Teradata estimate? **4** (800 / 200)
  5. How many people named 'Donaldson' would Teradata estimate? **16** (800 / 50)
  6. How many people named 'Cooper' would Teradata estimate? **6** (900 / 300)

**What to COLLECT STATISTICS On?**

You don't COLLECT STATISTICS on all columns and indexes because it takes up too much space for unnecessary reasons, but you do collect on:

- All Non-Unique Primary Indexes and All Non-Unique Secondary Indexes
- Non-indexed columns used in joins
- The Unique Primary Index of small tables (less than 1,000 rows per AMP)
- Columns that frequently appear in WHERE search conditions or in the WHERE clause of joins
- Primary Index of a Join Index
- Secondary Indexes defined on any join index
- Join index columns that frequently appear on any additional join index columns that frequently appear in WHERE search conditions

The first time you collect statistics, you collect them at the index or column level. After that, you just collect statistics at the table level and all previous columns collected previously are collected again. It's a mistake to collect statistics only once and then never do it again. COLLECT STATISTICS each time a table's data changes by 10%.

#### Why Collect Statistics?

- What does collecting statistics do to help the PE come up with a better plan?
- **Access Path** – The PE will easily choose and use any Primary Index access (UPI or NUPI), and it will also easily choose a Unique Secondary Index (USI), but statistics really help the PE decide whether or not to do a Full Scan or a Non-Unique Secondary Index (NUSI) or if it can use multiple NUSI's ANDed together to perform a NUSI bitmap.
  - **Join Method** – When you collect statistics, it gives Teradata a better idea whether or not to do a merge join, product join, hash join, or nested join.
  - **Join Geography** – When two rows are joined together, they must physically be located on the same AMP. The only way that this happens naturally is if the join column (PK/FK) is the Primary Index of both tables. Most of the time, this is not the case. When doing a join, the PE will analyze the paths of how it will read the rows from both tables on the same AMP. Will it redistribute (refresh by the partition column) over both of the tables, or will it duplicate the smaller table across all AMPS? A redistribution or duplication are the paths to co-location.
  - **Join Order** – All joins are performed two tables at a time. What will be the best order to join the tables together? When two or more tables are involved, this becomes very important.

It is the access path, the join method, the join geography, and the order that makes statistics collection so vital to all Teradata systems.

#### How do you know if Statistics were collected on a Table?

Syntax: HELP Statistics <Table Name>

Help Statistics Employee_Table ;			
Date	Time	Unique Values	Column Names
12/04/19	20:50:12	6	Employee_No
12/04/19	20:50:13	4	First_Name
12/04/19	20:50:130	2	MidName, First_Name



2

Help Statistics\_Hierarchy\_Table;

ERROR [HY000] [Teradata][ODBC Teradata Driver][Teradata Database]  
There are no statistics defined for the table.  
HELPSTATISTICS Command Failed.

Careful: This looks like an error,  
but it is merely stating that  
No Statistics were Collected!

The HELP Statistics command will show you what statistics have been collected, or specifically tell you that no statistics were collected on the table.

A Huge Hint that No Statistics Have Been Collected

EXPLAIN SELECT \* FROM New\_Employee\_Table ;

3) We do an all-AMPS RETRIEVE step from SQL CLASS New\_Employee\_Table by way of an all-rows scan with no residual conditions into Spool1 (group\_amps), which is built locally on the AMPS. The size of Spool 1 is estimated with low confidence to be 12 rows (64 bytes). The estimated time for this step is 0.03 seconds.

COLLECT\_STATISTICS ON New\_Employee\_Table  
column Employee\_Nr ;

EXPLAIN SELECT \* FROM New\_Employee\_Table ;

3) We do an all-AMPS RETRIEVE step from SQL CLASS New\_Employee\_Table by way of an all-rows scan with no residual conditions into Spool1 (group\_amps), which is built locally on the AMPS. The size of Spool 1 is estimated with high confidence to be 3 rows (513 bytes). The estimated time for this step is 0.03 seconds.

If you run an Explain on a query and the row estimate has No Confidence or Low Confidence, then that is a sign that no statistics were collected. Notice how the Explain above changed to High Confidence after we collected statistics on the table.



**The Basic Syntax for COLLECT STATISTICS**

- 1 COLLECT STATISTICS on <Tablename>  
COLUMN <Column Name> ;
- 2 COLLECT STATISTICS on <Tablename>  
INDEX (<Column Name(s)>);

Here are three actual examples:

- A** COLLECT STATISTICS on Employee\_Table  
COLUMN Employee\_No ;
- B** COLLECT STATISTICS on Employee\_Table  
COLUMN Dept\_No ;
- C** COLLECT STATISTICS on Employee\_Table  
INDEX (First\_Name, Last\_Name);

Employee\_No is the Primary  
Index. Always collect at the  
column level for single indexes.Only collect at the Index level  
for multi-column indexes.

The example commands above provide good fundamentals and concepts to follow.

**COLLECT STATISTICS Examples for a better Understanding**

- COLLECT STATISTICS on Employee\_Table  
COLUMN Dept\_No ;
- COLLECT STATISTICS on Employee\_Table  
COLUMN Employee\_No ;
- COLLECT STATISTICS on Employee\_Table  
INDEX(First\_Name, Last\_Name);
- COLLECT STATISTICS on Employee\_Table  
Column(Employee\_No, Dept\_No)
- COLLECT STATISTICS on Employee\_Table

Here is how you COLLECT  
STATISTICS on a single column.Here is how you COLLECT  
STATISTICS on a single index.Here is how to COLLECT  
STATISTICS on a Multi-column Index.Here is how to COLLECT  
STATISTICS on multiple columns  
that are often used together in the  
SQL WHERE Clause.Here is how to Refresh  
STATISTICS on columns and indexes previously  
collected or when the table has  
changed by at least 10%.**The New Teradata V14 Way to Collect Statistics**In previous versions, Teradata required that you had to Collect Statistics for each column separately, thus always  
performing a full table scan each time. Those days are over!

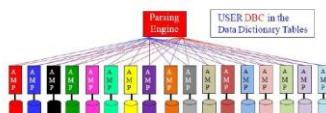
Old Way      New Teradata V14 Way



```
COLLECT STATISTICS COLUMNS First_Name,Last_Name  
ON Employee_Table;  
COLLECT STATISTICS COLUMNS First_Name,Last_Name  
ON Employee_Table;  
COLLECT STATISTICS COLUMNS Dept_No  
ON Employee_Table;  
COLLECT STATISTICS COLUMNS Dept_No  
ON Employee_Table;
```

The new way to collect statistics in Teradata V14 is to do it all at the same time. This is a much better strategy. Only a single table scan is required, instead of 3 table scans using the old approach. This is an incredible improvement.

Where Does Teradata Keep the Collected Statistics?



## Teradata CoE

## Teradata Lab Course

The Collect Statistics information is kept in user DBC in the Data Dictionary. The rows are spread evenly across all AMPs in three tables.

1. DBC **Indexer**(for multi-column indexes only)
2. DBC **TVFields**(for all columns and single column indexes)
3. DBC **StatsTbl**(Teradata V14 and beyond)

DBC is the most powerful user, and DBC owns the Data Dictionary (DD) so it makes sense that DBC will house the statistics in DBC tables. In V14, the statistics are housed in the DBC.StatsTbl relieving the contention for the DBC Indexes and TVFields tables.

### The Official Syntaxes for COLLECT STATISTICS

#### ① Syntax 1

```
COLLECT STATISTICS [ USING SAMPLE ]
[ OR [ TEMPORARY ] { <table-name> | <join-index-name> | <hash-index-name> }
| COUNT [ <column-name> | <column-list> ]
| [ UNIQUE ] INDEX [ <index-name> | ALL ] [ <column-list> ]
| [ UNIQUE ] INDEX [ <index-name> | ALL ] [ <column-name> ] ]
[ ORDER BY [ HASH | VALUES ] { <column-name> } ] ] ;
```

#### ② Syntax 2

```
COLLECT STATISTICS [ USING SAMPLE ]
[ COUNT [ <column-name> | <column-list> ]
| [ UNIQUE ] INDEX [ <index-name> | ALL ] [ <column-list> ]
| OS [ TEMPORARY ] { <table-name> | <join-index-name> | <hash-index-name> } ] ;
```

### How to Recollect STATISTICS on a Table

Here is the syntax for re-collecting statistics on a table

```
COLLECT STATISTICS ON <tablename> ;
```

Below is an actual example

```
COLLECT STATISTICS ON employee_Table;
```

This will NOT Collect on  
all columns in the table, but  
only refresh the columns and  
indexes currently collected on.

The first time you collect statistics, you do it for each individual column or index that you want to collect on. When a table changes its data by 10% due to Inserts, Updates, or Deletes, you merely use the command above, and it re-collects on the same columns and indexes previously collected on.

Teradata Always Does a Random AMP Sample



**Teradata CoE**      **Teradata Lab Course**

---

The Parsing Engine will hash the Table ID for a table being queried, and then use the Hash Map to determine which AMP will be assigned to do a Random AMP Sample for this table.

Remember that a Random AMP Sample only applies to **indexed columns** and **table row counts**.

In the old days, Teradata never did a Random AMP Sample unless statistics were not collected. But these days Teradata always does a Random AMP Sample before placing the Table Header inside each AMP's FSG Cache. This allows Teradata to compare these statistics with collected statistics to determine if statistics are old and stale. If the statistics are determined to be out of date, then the Random AMP Sample is used.

**Random Sample is kept in the Table Header in FSG Cache**

**Random Sample is kept in the Table Header in FSG Cache**

If the Sales\_Table Header is not in FSG Cache, I need for one AMP to do the random sample of the table's indexes, and then share it.

Then, check the Data Dictionary to see if statistics were collected.

If NO statistics were collected, then use the random sample.

If statistics were collected, then compare the random AMP sample to the collected statistics to determine if statistics are stale.

If statistics are stale, use the random sample.

Teradata compares the collected statistics to a Random AMP Sample (obtained by sampling a single AMP before placing the Table Header in FSG Cache). This compare determines if the statistics will be used or if they should be replaced by the sample.

**Multiple Random AMP Samplings**

The PE does Random AMP Sampling based on the Table ID. The Table ID is hashed and that AMP is always assigned the random sample for that table. This assumes that no single AMP will be tasked for too many tables, but if the table is badly skewed this can confuse the PE.

So now more than one AMP can be sampled when generating row counts for a query plan for much better estimations on row count, row size, and rows per value estimates per table.

In the DBS Control area field 65 can now set the standard for how AMPs are sampled.

65. **RandomAmpSampling** – this field determines the number of AMPs to be sampled for getting the row|

---

estimates of a table. The valid values are D, L, M, N or A.

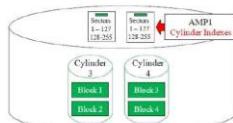
- D - The default is one AMP sampling (D is the default unless changed)
- L - Maximum of two AMP's sampling
- M - Maximum of five AMP's sampling
- N - Node Level sampling (all the AMPs in a node are sampled)
- A - System Level sampling (all the AMPs in the system are sampled)

Multiple AMP's can now be used for the random AMP sampling, so a higher number of AMP's sampled will provide better estimates to counter skewed results. But, it can cause short running queries to run slower just so long running queries can run faster.

#### How a Random AMP gets a Table Row count

For now count read 1 or 2 cylinders from 1 (or more) AMP's. Calculate the number of rows in the table by taking the:

1. Average Number of Rows per Block in the sampled Cylinder(s)
2. Multiply this by the Number of Data Blocks in the sampled Cylinder(s)
3. Multiply this by the Number of Cylinders for this table on this AMP(s)
4. Multiply this by the Number of AMP's in the system



If a table (or index subtable) spans more than 1 cylinder, it will sample the first and the last cylinder. If it fits into 1 cylinder, it will only sample that one cylinder. This way, the AMP can do some simple math to estimate the total row count for a table.

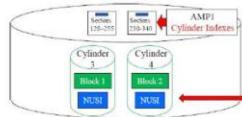
#### Random AMP Estimates for NUSI Secondary Indexes

For Non-Unique Secondary Indexes (NUSI) estimates, read 1 or 2 cylinders from the NUSI subtable, and then count the number of NUSI values in the cylinder(s).

The table row count is divided by the NUSI row counts to get a **Rows Per NUSI Value**.

It also assumes the number of distinct values on one AMP = total distinct values.





When you use a NUSI in the WHERE clause, the PE will often know if it needs to do a Full Table Scan. It is the statistics that help the PE in making this difficult decision. The Parser is more aggressive with COLLECTED STATISTICS. Features such as NTNU Bit Mapping require COLLECTED STATISTICS and no random sampling.

The Random AMP needs a couple of cylinders and some data NUSI blocks, and then does some simple math to estimate the Rows per NUSI Value. The PE then knows how strong or weak the WHERE clause is using the NUSI, and if it should even use the NUSI. This is the most important decision for the Parsing Engine. Should it just do a Full Table Scan, or use the NUSI? That is the biggest reason the PE needs statistics. That is why you should always collect statistics on all NUSI indexes.

#### USI Random AMP Samples are Not Considered

Random AMP sampling assumes that the number of distinct values in a USI needs to be small enough, so it does not need the USI suitable for USI equality conditions. Because equality conditions on a unique index return only one row, the PE automatically uses the USI without considering statistics. However, if a USI will frequently be specified in non-equality statements, such as range constraints, then you should collect statistics on the Unique Secondary Index.



You really only need to collect statistics on a Unique Secondary Index column if there are a lot of SQL statements on non-equality conditions such as range queries.

#### There's No Random AMP Estimate for Non-indexed Columns

For non-indexed columns without statistics, the optimizer uses a fixed formula to estimate the number of rows. This is sometimes referred to as **heuristics**.

Teradata assumes 10% for one column in an equality condition in the WHERE clause.

Assumes 7.5% for two columns, each in an equality condition, and ANDed together.



## Teradata CoE

## Teradata Lab Course

```
SELECT *  
FROM Sales_Table  
WHERE Product_ID = 1000;
```

10% of the rows have  
Product\_ID of 1000

```
SELECT *  
FROM Sales_Table  
WHERE Product_ID = 1000  
AND Daily_Sales = 50000 ;
```

7.5% of the rows will qualify  
because of having a  
Product\_ID of 1000 and  
a Daily\_Sales of 50000

Teradata does not do a Random AMP Sample for non-indexed columns that are used in the WHERE clause of the SQL. It uses the above for a quick and dirty estimate.

### Summary of the PE Plan if No Statistics Were Collected

Today's Teradata systems always perform a random AMP sample even if tables have statistics. Then, they compare the random AMP sample with the statistics to determine if the statistics are stale.

A random AMP is selected for a random sample. Two things happen:

- 1) Indexes are sampled on the random AMP, and the PE estimates based on the total number of AMPS in the system.
- 2) If a column in the WHERE clause of the SQL is not an index, the PE assumes that 10% of the rows will come back. If two columns are in the WHERE clause, it assumes 5%. If three columns are in the WHERE Clause, it assumes 3%.

A random AMP sample is selected by the PE, so if it finds there are no statistics on the table or if the statistics are old and stale it has options.

### Stale Statistics Detection and Extrapolation

The PE estimates the table row count based on the Primary Index collection of a table (called the histogram), but it also does a Random AMP Sample for comparison. If these two metrics differ by a threshold of 10% or for small tables (10,000 new rows), then statistics are considered stale.

Any cardinality estimations will now use extrapolation. What is extrapolation? This means that the PE will estimate based on past statistics and go with new estimations. In other words, the PE will derive its own estimate based on past history and the random sample.

Small tables with less than 25 rows per AMP, or for skewed tables, no extrapolation will done.

Primary Index statistics 150,000	Parsing Engine
	<ol style="list-style-type: none"><li>1. The data isn't skewed.</li><li>2. The threshold is &gt; 10%</li><li>3. The table is not small.</li><li>4. I think the value is 500,000!</li></ol>
Random AMP Sample 500,000	

When statistics are determined to be stale, the PE will use the Random AMP sample and also extrapolate, which means to estimate new statistics based on historical data.

### Extrapolation for Future Dates



AMP 1	AMP 2
Order_Table	Order_Table
Jan 2014	Jan 2014
Feb 2014	Feb 2014
Mar 2014	Mar 2014
Apr 2014	Apr 2014
May 2014	May 2014
Jun 2014	Jun 2014
Jul 2014	Jul 2014
Aug 2014	Aug 2014
Sep 2014	Sep 2014
Oct 2014	Oct 2014
Nov 2014	Nov 2014
Dec 2014	Dec 2014

```
EXPLAIN
SELECT *
FROM Order_Table
WHERE Order_Date BETWEEN
'2015-01-01' AND '2015-03-31';
```

**Parsing Logins**

1. Only have statistics for 2014.
2. The first quarter of 2014 had 1,000,000 orders.
3. Orders have been increasing by 10%.
4. I think I can extrapolate and estimate that there will be 1,100,000 orders in the first quarter of 2015.

Above, you can see the PC only has collect statistics for the year 2014, but the Explain is asking about 2015 data. Teradata will extrapolate new estimates for future dates based on history and growth estimates. This allows for better estimates and less concern about statistics collection.

**How to Copy a Table with Data and the Statistics?**

This next example is pretty amazing. Assume that the original Employee\_Table had COLLECT STATISTICS on the columns Employee\_No and Dept\_No. The new table we have called Employee\_New\_Table. New will have DDL exactly like the Employee\_Table plus data plus the statistics. Yes, the exact same statistics will be copied to the new table. Below is the actual example:

```
CREATE TABLE Employee_Table_New AS Employee_Table
WITH DATA
AND STATISTICS;
```



The example above will CREATE a new table called Employee\_Table\_New, and it will have the exact same DDL as the Employee\_Table, the exact same data, and the exact same statistics.

**COLLECT STATISTICS Directly From another Table**

## Teradata CoE

## Teradata Lab Course

CREATE TABLE Stats\_Test AS  
(Select \* from Employee\_Table)  
WITH DATA AND STATISTICS;

Data and Statistics  
have been copied

CREATE TABLE Stats\_Test2 AS  
(Select \* from Employee\_Table)  
WITH DATA;

Experts copy the  
Statistics!

COLLECT STATISTICS ON Stats\_Test2  
FROM Employee\_Table;

Statistics  
copied!

In Teradata V13 and above, you can Collect Statistics directly from another table.

How to Copy a Table with NO Data and the Statistics?

This next example is clever. Assume that the original Employee\_Table had COLLECT  
STATISTICS on the columns Employee\_No and Dept\_No. The new table we have called  
Employee\_Table\_99 will have DDL exactly like the Employee\_Table but  
NO data. It will have the Statistics, but they will be **Zeroed** Statistics.

CREATE TABLE Employee\_Table\_99 AS Employee\_Table  
WITH NO DATA;  
AND Statistics;

Table DDL copied with NO Data and Zeroed Statistics

on the columns Employee\_No and Dept\_No.

Once you have loaded Employee\_Table\_99 with the data you want by using  
BTEQ, FastLoad, an INSERT-SELECT, Nexus, or the ETL tool of your  
choice so you can then ReCollect Statistics!

COLLECT STATISTICS ON Employee\_Table\_99;

You have just Re- Collected Statistics on Employee\_Table\_99 for the columns Employee\_No and Dept\_No. The  
key re-collecting the statistics is that when you **load** the data is loaded for the purpose of  
getting the Zeroed Statistics in the first place. Make sure you **recollect** after your data is loaded though!

When to COLLECT STATISTICS Using only a SAMPLE

You might consider Collecting Statistics with SAMPLE if  
You are collecting statistics on a very **large** table.

When collecting statistics becomes a problem with system performance or cost because the **system is so**

**slow**



Don't consider Collecting Statistics with SAMPLE if:

- 1 The tables are **small**.
- 2 To **replace** all existing full scan Collected Statistics.

- 3 If the column's data is **skewed badly**.

COLLECT STATISTICS can be very time consuming because it performs a full table scan and then performs a lot of statistical calculations. Because Collect Statistics runs infrequently and benefits query optimization, it is considered a necessary task. Without statistics, query performance will suffer. The bad news about sampled statistics is that they may not always provide the best statistics for the PEs. This may affect the PE's plans. In most cases, sampled statistics are better than no statistics. Don't use Sample unless necessary.

Examples of COLLECT STATISTICS Using only a SAMPLE

COLLECT STATISTICS USING SAMPLE on Employee\_Table COLUMN Dept\_Nr ;

COLLECT STATISTICS USING SAMPLE on Employee\_Table COLUMN Employee\_Nr ;

COLLECT STATISTICS USING SAMPLE on Employee\_Table INDEX (First\_Name\_Last\_Name);

COLLECT STATISTICS on Employee\_Table ;

Here is how you COLLECT STATISTICS on a single column.

Here is how you COLLECT STATISTICS on a single index.

Here is how to COLLECT STATISTICS on a Multi-column Index.

Here is how to Refresh STATISTICS on a Table.

If you **recollect** statistics on a **sample**, it recollects with the **same** sample amount!

Sampled statistics are generally more accurate for data that is not skewed. For example, columns or indexes that are unique or nearly unique are not skewed. Because the PE needs to be aware of skewed data, you should not collect with sample on skewed data. That is why sampling is generally more appropriate for indexes than non-indexed column(s). If you recollect statistics on a Sample, it Recollects with the same Sample!

Examples of COLLECT STATISTICS for V14



To collect sample statistics using the [system default sample](#):

```
COLLECT STATISTICS USING SYSTEM SAMPLE COLUMN (Product_ID) ON Sales_Table;
```

To collect sample statistics by scanning [15 percent](#) of the rC1Ws and use [100 intervals](#):

```
COLLECT STATISTICS USING SAMPLE 15 PERCENT AND MAXINTERVALS 100
```

COLLECT STATISTICS USING SAMPLE 15 PERCENT AND MAXINTERVALS 100

COLUMN (Product\_ID) AS Product\_Stats ON Sales\_Table;

To change sample statistics to [20 percent](#) (or Product\_ID) and use [250 intervals](#):

```
COLLECT STATISTICS USING SAMPLE 20 PERCENT AND MAXINTERVALS 250
```

COLUMN (Product\_ID) AS Prod\_stats ON Sales\_Table;

To display the COLLECT STATISTICS statements for a table:

```
SHOW STATISTICS ON Sales_Table;
```

To display statistics [details](#) - summary section, high bias values, and intervals:

```
SHOW STATISTICS VALUES COLUMN Product_ID ON Sales_Table;
```

#### How to Collect Statistics on a PPI Table on the Partition

Here is the syntax for collecting statistics on a PPI table.

```
COLLECT STATISTICS on <Tablename> COLUMN <Column> PARTITION
```

Here is an actual example:

```
COLLECT STATISTICS on Order_Table_PPI COLUMN Partition;
```

Three reasons to Collect on the Partition:

The Parsing Engine will have a better plan for PPI Tables.

This is especially useful with Partition Elimination on Range Queries.

This is especially helpful when a table has a lot of empty partitions.

The Parsing Engine can use this information to better estimate the query cost when there are a significant number of empty partitions. If PARTITION statistics are not collected, empty partitions may cause the Parsing Engine to underestimate the number of rows in a partition. You shouldn't use WITH SAMPLE to collect on Partitions.

Teradata V12 and V13 Statistics Enhancements

In V12 Extrapolate Statistics is designed to more accurately provide for a statistical estimate for date range-based queries that specify a "future" date that is outside the bounds of the current statistics. This results in less biased collections.

In V12 State Statistics Detection compares the Random AMP Sample with the statistics collected and determines if they are stale and should not be used.

In V13 Statistics can now be collected on Volatile Tables.

In V13 PARTITION statistic capabilities have been added to Global Temporary Tables. In H43 Multi-Column statistics are now available on User and Job Indexes.

Indexes.

In V13 Sample Statistics are available on Tables, Volatile Tables, Global Temporary Tables, Hash Indexes and Join Indexes, including the Partition Columns.

#### Teradata V14 Statistics Enhancements

- There is now a **SUMMARY** option to collect table-level statistics.
- SYSTEM SAMPLE** option allows the system to determine the sampled system percentage.
- Sampling options have been enhanced (e.g., SAMPLE n PERCENT).
- Statistics are stored in **DBC.StatsTbl** to reduce access contention and improve performance.
- New views DBC.StatsV, DBC.ColumnStatsV, DBC.MulticolumnStatsV, and IndexStatsV.
- SHOW STATISTICS statement reports detailed statistics in plain text or XML formatting.
- Internal PE enhancements for histogram structure and use, including:
  - Storing statistics data in their native Teradata data types without losing precision
  - Enhanced extrapolation methods for stale statistics
  - Maintaining statistics history

Teradata V14 now allows you to determine a sampling percentage for sampled statistics. You can even collect/recollect either all or a portion of the data in a table. You can also collect stats on global temporary tables, and you can provide a name for statistics collected while also being able to specify the column ordering for multicolumn statistics. There is also a dedicated statistics cache that is designed to improve query optimization time.

#### Teradata V14 Summary Statistics

New in Teradata 14.0, table-level statistics known as "summary statistics" are collected whenever column or index statistics are collected. Summary statistics do not cause their own histogram to be built, but rather they create histograms for the columns undergoing collection that are held in the new **DBC.StatsTbl**. Here are some of the items in "summary statistics":

- Row count
- Average block size
- Block level compression metrics
- Temperature

One critical advantage is that the optimizer now uses summary stats to get the most up-to-date row count from the table in order to provide more accurate extrapolations.

Here is how you can see the Summary Statistics:

```
SHOW SUMMARY STATISTICS VALUES ON Employee_Table;
```

You can specifically request summary statistics for a table, but you never need to do that because each individual statistics collection statement causes summary stats to be gathered (it is a quick process). It is best in V14 to now write your collection scripts to do it all in one scan Vs. multiple statements.

#### Teradata V14 MaxValueLength

```
COLLECT STATISTICS  
USING MAXVALUELength 50  
ON ProductTable  
OFF Product_Table /
```



## Teradata CoE

## Teradata Lab Course

Before V14, whenever you collected statistics, Teradata only placed the first 16 bytes in the statistics so long names were cut off. Now, the default is 25 bytes, but you can use the MaxValueLength keyword (example above) to specify the length you want.

MAXVALUELength lets you expand the length of the values contained in the histogram for that statistic. The new default length is 25 bytes, when previously it was 16. If needed, you can specify well over 1000 bytes for a maximum value length. The 16-byte limit on value sizes in earlier releases was always padded to 16 bytes. Now, the length can be longer, but no padding is done.

### Teradata V14 MaxIntervals

```
COLLECT STATISTICS
    FOR Employee_Table
    COLUMNS (Last_Name)
    ON Employee_Table;
```

Before V14, whenever you collected statistics, Teradata did a full table scan on the values, sorted them, and then placed them into 200 intervals. Now, the default is 200 intervals, but you can specify (example above) the number of intervals you desire.

Each statistics interval highlights its single most popular value, and the number of rows that carry that value are recorded. The rest of the values in the interval are ignored. By reducing the number of intervals, the optimizer can accurately get valuable information about the distribution of the most popular values. It is useful if most of your data is highly widespread skew on a column or index you are collecting statistics on and you want more individual high-row-count values to be represented in the histogram. The range is 0 - 500 for MaxIntervals.

### Teradata V14 Sample N Percent

```
COLLECT STATISTICS
    FOR Employee_Table
    COLUMNS (Last_Name)
    ON Employee_Table;
```

Using Sample before defaulted each time to the system parameter, but now you can specifically state the percent you want for each column or index.

SAMPLE n PERCENT allows you to specify sampling at the individual statistics collection level, rather than at the system level. This allows for individual statistics sampling to different columns and indexes can be performed. The better you get at knowing your data and the queries upon them, the more you can specifically use the sampling to better help the PE.

### Teradata Statistics Wizard

The Teradata Statistics Wizard is a graphical tool that automates the collection and re-collection of statistics, helping the DBA to better manage the statistics process.

Now the DBA can specify a workload to be analyzed and receive statistics recommendations.

Any database or collection of tables, indexes, or columns can be selected for the collection, or re-collection of statistics.

Recommendations can be based on table demographics and general heuristics.

Define execution of and schedule the arbitrary collection/re-collection for later.

Display and modify the interval statistics for a column or index.

Provides numerous reports on statistics recommendations, update cost analysis and table usage.

The Statistics Wizard can be used to help the DBA with statistics.



## Data Manipulation Language



**Data Manipulation Language (DML)**



**INSERT Syntax #1**

The following syntax of the INSERT does not use the column names as part of the command. Therefore, it requires that the VALUES portion of the INSERT match each and every column in the table with a data value or a NULL.

```
[INT|INT2|INT4] <table-name> VALUE  
[<literal>-<data-value>|...<literal>-<data-value>|...]
```

Note: Using INS instead of INSERT is not ANSI compliant.

The INSERT statement is used to put a new row into a table. A status is the only returned value from the database, no rows are returned to your. It must account for all the columns in a table using either a data value or a NULL. When executing the INSERT places a single new row into a table. Although it can run as a single row insert, primarily it is used in utilities like BTEQ, FastLoad, MultiLoad, Trimp or some other application that reads a data record and uses the data to build a new row in a table.

**INSERT Example with Syntax 1**

```
INSERT INTO My_Table VALUES  
(‘My character data’,124.56, 102587, ,NULL, ‘2000-12-31’);
```



My character data	124.56	102587	NULL	NULL	2000-12-31
-------------------	--------	--------	------	------	------------

After the execution of the above INSERT, there is a new row with the first character data value of ‘My character data’ going into Column1, the decimal value of 124.56 into Column2, the integer 102587 into Column3, NULL values into Column4 and Column5, and a date into Column6.

The NULL enclosed in the VALUES list is the literal representation for no data. However, the two commas (,) that follow the NULL value in the Column4 also represent missing data. The commas are placeholders or delimiters for the data values. When no data value is coded, the end result is a NULL.

**INSERT Syntax #2**

The syntax of the second type of INSERT follows:

```
[INT|INT2|INT4] <table-name>  
<column-name>|,<column-name>|  
[VALUES  
[<literal>-<data-value>|...<literal>-<data-value>|...]]
```

Note: Using INS instead of INSERT is not ANSI compliant.

This is another form of the INSERT statement that can be used when some of the data is not available. It allows for the missing values (NULL) to be eliminated from the list in the VALUES clause. It is also the best format when the data is arranged in a different sequence than the CREATE TABLE, or when there are more nulls (unknown values) than available



data values.

#### INSERT Example with Syntax 2

```
INSERT INTO My_Table (Column2, Column1, Column3, Column6)
VALUES ( 124.56, 'My character data', 12587, 2000-12-31);
```



My character data	124.56	102587	NULL	NULL	2000-12-31
-------------------	--------	--------	------	------	------------

The above statement incorporates both of the reasons to use this syntax. First, notice that the column names Column2 and Column1 have been switched, to match the data values. Also, notice that Column1 and Column5 do not appear in the column list, therefore they are assumed to be NULL. This is a good format to use when the data is coming from a file and does not match the order of the table columns.

#### INSERT Example with Syntax 3

```
INSERT INTO My_Table
(Column2=124.56, Column1='My character data', Column3=12587,
Column6='2000-12-31');
```



My character data	124.56	102587	NULL	NULL	2000-12-31
-------------------	--------	--------	------	------	------------

The third form of the INSERT statement can be used to insert the same row as the previous INSERT. It might look like this.

#### Using NULL for Default Values

Either of the next two INSERT statements may be used to build a row with no data values in My\_Table

```
INSERT INTO My_Table VALUES (.....);
```

```
INSERT INTO My_Table VALUES
(NULL,NULL,NULL,NULL,NULL,NULL);
```



NULL	NULL	NULL	NULL	NULL	NULL
------	------	------	------	------	------

Teradata now has the ANSI DEFAULT VALUES functionality. Although an INSERT statement could easily put a null value into a table column, it requires it to use the NULL reserved word or by omitting a value for that column(s) between commas.

#### INSERT/SELECT Command

The syntax of the INSERT / SELECT is:

```
[INS[ERT] [INTO] <table-name>
```



```
SELECT <column-name>, ...<column-name>
```

Although the INSERT is great for adding a single row not currently present in the system, an INSERT/SELECT is even better when the data already exists within Teradata. In this case, the INSERT is combined with a SELECT. However, no rows are returned to the user. Instead, they go into the table as new rows.

The SELECT reads the data values from the one or more columns in one or more tables and uses them as the values to INSERT into another table. Simply put, the SELECT takes the place of the VALUES portion of the INSERT.

This is a common technique for building data marts, interim tables and temporary tables. It is normally a better and much faster alternative than extracting the rows to a data file, then reading the data file and inserting the rows using a utility.

#### INSERT/SELECT Example using All Columns (\*)

When all columns are desired to make an exact copy of the second table, and both tables have the exact same number of columns in the exact same order with the exact same data types. An \* may be used in the SELECT to read all columns without a WHERE clause, as in the example below:

```
INSERT INTO My_Table
SELECT * FROM My_Original_Table;
```

Like all SELECT operations without a WHERE clause, a full table scan occurs and all the rows of the second table are inserted into My\_Table using only the data values from the columns listed.

#### INSERT/SELECT Example with Less Columns

When fewer than all the columns are desired, either of the following INSERT / SELECT statements will do the job:

```
INSERT INTO My_Table
SELECT Column1, Column2, Column3, ...
FROM My_Original_Table;

INSERT INTO My_Table (Column2, Column1, Column3, Column8)
SELECT Column2, Column1, Column3, CURRENT_DATE
FROM My_Original_Table;
```

In both of the above examples, only the first three and the last columns are receiving data. In the first INSERT, the data is a literal date. The second INSERT uses the CURRENT\_DATE. Both are acceptable, depending on what is needed.

Working with the same concept of a normal INSERT, when using the column names, the only data values needed are for those columns. It must be in the same sequence as the column list, not the CREATE TABLE. Therefore, omitted data values or column names become a NULL data value.

#### INSERT/SELECT to Build a Data Mart

As an example of a data mart, it might be desirable to build a summary table using something like the following:

```
INSERT INTO My_Table
SELECT 1 AS ID, SUM(column1), AVG(column3),
       MAX(column4), AVG(CAST(column5)),
       AVG(CAST(column6)) 
  FROM My_Sales_Table
 GROUP BY 1;
```

However used, the INSERT / SELECT can be a powerful tool for creating rows from the rows already contained in one or more other tables.

#### Fast Path INSERT/SELECT



```
INSERT INTO My_Table SELECT * FROM My_Original_Table ;j
```

When the table being loaded is empty, the INSERT / SELECT is very fast. This is especially true when all columns and all rows are being copied. Remember, the table being loaded must be empty to attain the speed. If there is even one row already in the table, it negates the ability to take the Fast Path.

This occurs because the rows in the spool. Prior to the insert, Teradata needs to "Transient" Journal an identifier for each inserted row. Recovery, if needed, is to ensure the table has no other type of recovery can be easier or faster.

When all columns and all rows are requested from the existing table, and they exactly match the columns in the new table, there is no need to use spool. The rows go straight into the table being loaded. Additionally, when all rows are being selected, Teradata does not bother to read the individual rows. Instead, each AMP literally copies the blocks of the original table to those for the new table.

These actions all unify it is called the Fast Path. To use this technique, the order of the columns in both tables must match exactly, or else the data types. Otherwise, spool must be used to rearrange the data values or translate from one data type to another.

#### NOT quite the Fast Path INSERT/SELECT

What if it is necessary to retrieve the rows from multiple tables for the INSERT?

Multiple INSERT / SELECT operations could be performed as follows:

```
INSERT INTO My_Table    SELECT * FROM My_Original_Table_1;
INSERT INTO My_Table    SELECT * FROM My_Original_Table_2;
INSERT INTO My_Table    SELECT * FROM My_Original_Table_3;
```

The first SELECT into My\_Table loads the table quickly for even with millions of rows. However, the table is no longer empty, and the subsequent INSERT is much slower because it cannot use the fast path. All inserted rows must be identified in the Transient Journal. It can more than double the processing time.

The real question is: How does one make the entire individual SELECT operations act as one, so that the table stays empty until all rows are available for the INSERT?

#### UNION for the Fast Path INSERT/SELECT

One way get the Fast Path is to use the UNION command to perform all SELECT operations in parallel before the first row is inserted into the new table. Therefore, all rows are read from the various tables combined into a single answer set in spool, and then loaded into the empty table. All of this is done at high speed.

For instance, if all the rows from three different tables are needed to populate the new table, the applicable statement might look like the following:

```
INSERT INTO My_Table
SELECT * FROM My_Original_Table_1
UNION
SELECT * FROM My_Original_Table_2
UNION
SELECT * FROM My_Original_Table_3;
```

Again, the above statement assumes that all four tables have exactly the same columns. Whether or not that would ever be the case in real life, this is used as an example. However, at this point we know the columns in the SELECT must match the columns in the table to be loaded, no matter what that is accomplished.

#### BTEQ for the Fast Path INSERT/SELECT

A second alternative method is available using BTEQ. The key here is that BTEQ can do multiple SQL statements as a single transaction. The SELECT and the INSERT operations. The only way to do that is to delay the actual INSERT, until all of the rows from the selected operations have completed. Then, the INSERT is performed as a part of the same transaction into the empty table.



```
.logon local1/dbc  
Password:*****  
Session successfully completed  
BTEQ - Enter your BTEQ/SQL request or BTEQ command:  
INSERT INTO My_Table  
SELECT * FROM My_Original_Table_1  
; INSERT INTO My_TableOriginal_Table_1  
SELECT * FROM My_Original_Table_2  
; INSERT INTO My_TableOriginal_Table_2  
SELECT * FROM My_Original_Table_3
```

By having another SQL command on the same line as the semi-colon (;), in BTEQ, they all become part of the same multi-statement transaction. Therefore, all are inserting into an empty table, and it is much faster than doing each INSERT individually.

#### The UPDATE Command Basic Syntax

```
UPDATE [<table-name>  
[FROM <table-name> [<alias-name>]]  
SET <column-name> = <expression>|<data-value> | <data-value> )  
<column-name> = <expression>|<data-value> | <data-value> ]  
[WHERE <condition-test>  
[AND <condition-test>] [ OR <condition-test> ] [ALL]
```

The UPDATE statement changes the data in one or more columns of one or more existing rows. A status is the only returned value from the database, no rows are returned to the user.

When business requirements call for a change to be made in the existing data, the UPDATE is the SQL statement to use. In order for the UPDATE to work, it must know a few things about the data row(s) involved. Like all SQL, it must know which table to use for making the change, which column or columns to change, and the like to make within the data.

#### Two UPDATE Examples



```
UPDATE My_Table
SET Column1 = 256
    ,Column2 = 'None'
    ,Column3 = 'Yours'
WHERE Column4 = 'My character data'.
```

```
UPDATE My_Table
SET Column1 = Column2 + 256
WHERE Column1 = 'My character data'
    ,Column2 = 'None'
    ,Column3 = 'Yours';
```

The first UPDATE command modifies all rows that contain 'My character data' including the one that was inserted earlier in this chapter. It changes the values in three columns with new data values provided after the equal sign (=).

The next UPDATE uses the same table as the above statement. However, this time it modifies the value in a column based on its current value and adds 256 to it. The UPDATE determines which row(s) to modify with compound conditions written in the WHERE clause based on values stored in other columns.

#### Subquery UPDATE Command Syntax

```
UPDATE[ table-name
      [ AS alias-name ] ] [ column-name ] = expression | data-value ]
      SET column-name = expression | data-value | <data-expression>|<data-value>
      WHERE column-name [ , column-name ] = expression | data-value | <data-expression>|<data-value>
            [ AND column-name [ AS alias-name ] = expression | data-value ]
                  FROM table-name [ AS alias-name ]
                  [ WHERE condition-text; ] ) [ ALL ] ;
```

Sometimes it is necessary to update rows in a table when they match rows in another table. To accomplish this, the tables must have one or more columns in the same domain. The matching process then involves either a subquery or join processing.

Notice in the above syntax example, that creating an alias for the table being updated is not compatible with a FROM clause. This change was made in V2R4.

#### Example of Subquery UPDATE Command

To change rows in My\_Table using another table called Ctl\_Tbl, the following UPDATE uses a subquery operation to accomplish the operation:

```
UPDATE My_Table
SET Column1 = 500000
WHERE Column1 IS ( SELECT Column1 FROM Ctl_Tbl
                    WHERE Column1 > 5000
                    AND Ctl_Tbl.Column1 IS NOT NULL );
```

Sometimes it is necessary to update rows in a table when they match rows in another table. To accomplish this, the tables must have one or more columns in the same domain. The matching process then involves either a subquery or join processing.

Notice above, that creating an alias for the table being updated is not compatible with a FROM clause, and remember that the change was made in V2R4.

#### Join UPDATE Command Syntax

```
UPDATE[ table-name
      [ AS alias-name ] ] [ column-name ] = expression | data-value
      SET column-name = expression | data-value | <data-expression>|<data-value>
      WHERE [ table-name ].column-name = [ table-name ].column-name
            [ AND condition-text; ] OR condition-text; ] ) [ ALL ] ;
```



When adding an alias to the UPDATE, the alias becomes the table name and MUST be used in the WHERE clause when qualifying columns.

#### Example of an UPDATE Join Command

To change rows in My\_Table using another table called Ctl\_Tbl, the following UPDATE uses a join to accomplish the operation. Both examples are equivalent:

```
UPDATE My_Table
  FROM Ctl_Tbl AS Ctbl
  SET Column1 = 2000000
    WHERE Column1 = 'A'
      WHERE My_Table.Column1 = Ctbl.Column2
        AND My_Table.Column3 > 5000 AND Ctl_Tbl.Column4 IS NOT NULL ;
```

In reality, the FROM is optional. This is because Teradata can dynamically include a table by qualifying the join column with the table name. The FROM is only needed to make an alias for the join table. The second UPDATE is the same as the above without the FROM for Ctl\_Tbl.

#### Fast Path UPDATE

The following INSERT/SELECT "updates" the values in Column3 and Column5 in every row of My\_Table, using My\_Table\_Copy

```
INSERT INTO My_Table_Copy
  SELECT Column1
    ,Column2
    ,Column3+1.05
    ,Column4
    ,Column5
  FROM My_Table ;
```

When the above command finishes, My\_Table\_Copy contains every row from My\_Table with the needed update.

#### The DELETE Command Basic Syntax

```
[DELETE] [ FROM ] <table> [ AS <aliasname> ] [ WHERE condition ] [ alias ]
```

The DELETE statement has one function, and that is to remove rows from a table. A status is the only returned value from the database, no rows are returned to the user. One of the fastest things that Teradata does is to remove ALL rows from a table.

The reason for its speed is that it simply moves all of the sectors allocated to the table onto the free sector list in the AMP's Cylinder Index. It is the fast path, and there is no COMMIT required unless the explicit transaction has not yet completed. In the event of a power failure, the rows are lost and the database must be restored before a COMMIT. Otherwise, the rows are gone and it will take either a backup tape or a BEFORE image in the Permanent Journal to perform a manual rollback. Be Very CAREFUL with DELETE. It can come back to bite you if you're not careful.

#### Two DELETE Examples to DELETE ALL Rows in a Table



## Teradata CoE

## Teradata Lab Course

DELETE FROM My\_Table [ ]

[ DEL My\_Table ]

Both examples will delete all the rows in the table.

Since the FROM and the ALL are optional, and the DELETE can be abbreviated, the second example still removes all rows from a table and executes exactly the same as the above statement.

A DELETE Example Deleting only Some of the Rows

DELETE FROM My\_Table [ ]

[ WHERE column < 100000 ]

The DELETE example above only removes the rows that contained a date value less than 1001231 (December 31, 2000).

in Column6 (DATE data type) and leaves all rows newer than or equal to the date.

Subquery and Join DELETE Command Syntax

The basic syntax for the DELETE statement

[DELETE] [table-name]  
[FROM] [table-name] [ ]  
[WHERE] [column-name] < condition ]  
[OR] [column-name] ( , , column-name )  
[AND] [column-name] ( , , column-name )  
[WHERE] [condition] [ ]

The join syntax for DELETE statement

[JOIN] [table-name]

[FROM] [table-name] [ AS alias ]  
[WHERE] [alias.]columnname=alias.columnname  
[AND] [condition] [ ]

Sometimes it is desirable to delete rows from one table based on their existence in or by matching a value stored in another table. For example, you may be asked to give a raise to all people in the Awards Table. To access these rows from another table for comparison, a subquery or a join operation can be used.

Example of Subquery DELETE Command

To remove rows from My\_Table using another table called Control\_Del\_Tbl, the next DELETE uses a subquery operation to accomplish the operation.

DELETE FROM My\_Table  
[ WHERE Column1, cd IN ( SELECT Column2 FROM Control\_Del\_Tbl ) ]

The above uses a Subquery and the DELETE command.

Example of Join DELETE Command

To remove the same rows from My\_Table using a join with the table called Control\_Del\_Tbl, the following is another technique to accomplish the same DELETE operation as the subquery example on the previous page.

DELETE FROM My\_Table  
[ JOIN ] [Control\_Del\_Tbl AS ch1, This AS ch2 ]  
[ WHERE My\_Table.Column1 = ch1.Column1  
AND ch1.Column2 = ch2.Column2 ]

This previous example could also be written using the formal JOIN clause. However, an alias cannot be used with this JOIN clause.

 Capgemini  
INNOVATION DRIVEN ENTERPRISE

```
DELETE My_Table
WHERE My_Table.Column2 = Control_Del.Th1.Column1
    AND My_Table.Column1 = Control_Del.Th1.Column1
    AND Control_Del.Th1.Column1 IS NULL;
```

The above uses a Join and the DELETE, and is equivalent to the previous subquery example we saw on the previous page.

#### Fast Path DELETE

Fast Path DELETE always occur when the WHERE clause is omitted.

However, most of the time, it is not desirable to delete all of the rows. Instead, it is more practical to remove older rows to make room for newer rows, or periodically purge data rows beyond the scope of business requirements.

For instance, the table is supposed to contain twelve months' worth of data, and it is now month number thirteen. It is now time to get rid of rows that are older than twelve months.

As soon as a WHERE clause is used in a DELETE, it must take the slow path to delete the rows. This simply means that it must log or journal a copy of each deleted row. This is to cover the potential that the command might fail. If that should happen, the system can roll back the transaction by performing a ROLLBACK. As slow as this additional processing makes the command, it is necessary to insure data integrity.

To use the Fast Path, a technique is needed that eliminates the journal logging. The trick is, again, to use a Fast Path INSERT / SELECT. This means, we insert the rows that need to be kept into an empty table.

#### Fast Path DELETE Example #1

Normal Path Processing for the DELETE (uses the Transient Journal):

```
DELETE FROM My_Table
WHERE Column6 < 1001231 ;
```

There are three different methods for using Fast Path Processing in BTEQ for a DELETE. The first method uses an INSERT / SELECT. It will be fast, but it will require a temporary table and appropriate DDL. It also requires that additional PERM space be available for temporarily holding both the rows to be kept and all of the original rows at the same time.

```
INSERT INTO My_Table_Copy
SELECT * FROM My_Table WHERE Column6 > 1001230
; DROP TABLE My_Table_Copy ;
; RENAME My_Table_Copy to My_Table ;
```

The first example does NOT use the Fast Path Delete, but the second example does.



**Fast Path DELETE Example # 2**

Normal Path Processing for the DELETE (uses the Transient Journal)

```
DELETE FROM My_Table
WHERE Column1 = '1001231' ;
```

This next method also uses an INSERT/SELECT and will be fast if it does: not require privileges for using any DCL. It probably will not be faster than the first method, since the rows must all be put back into the original table. However, the table is empty and the Fast Path will be used.

```
INSERT INTO My_Table_Copy
SELECT * FROM My_Table WHERE Column1 >= 1001230 ;
DELETE My_Table ;
CREATE GLOBAL TEMPORARY TABLE My_Table
SELECT * FROM My_Table_Copy ;
```

The first example does NOT use the Fast Path Delete, but the second example does.

**Fast Path DELETE Example # 3**

This INSERT/SELECT uses a Global Temporary table to prepare for the single transaction to copy My\_Table in

BTEQ

```
INSERT INTO My_Global_Table_Copy
SELECT * FROM My_Table WHERE Column1 >= 1001230 ;
DELETE My_Table ;
CREATE GLOBAL TEMPORARY TABLE My_Table
SELECT * FROM My_Global_Table_Copy ;
```

A Global Temporary Tables requires that TEMPORARY space be available for temporarily holding the rows to be kept and all of the original rows at the same time. A Volatile Temporary table could also be used. Its space comes from spool. However, it requires a CREATE statement to build the table, unlike Global Temporary tables. More information on Temporary tables is available in this book.

If you are not using BTEQ, these statements can be used in a macro. This works because macros always execute as a transaction.

**MERGE INTO**

Here are the V2RS MERGE Rules:

1. The Source relation must be a single row.
2. The Primary index of the target relation must use an equality condition to a numeric constant or a string literal in the ON Clause.
3. You cannot request an error log.

Here are the V12 and above MERGE Rules:

1. The Source relation can be either a single row or multiple rows.
2. The Primary index of the target relation must be bound by use of an equality condition to an explicit term or to a column set of the Source relation.
3. The Primary index of the target table cannot be updated.
4. The INSERT if the WHEN NOT MATCHED Clause is specified, must have the value specified in the ON clause with the target table primary index of the target table, which causes the INSERT to be on the local AMP.



## Teradata CoE

## Teradata Lab Course

MERGE merges a source row set into a target table based on whether there is a MATCH or whether there is a NOT MATCH condition. If there is a MATCH, then Teradata will UPDATE the row, but if there is a NOT MATCH condition, it will INSERT the row. This is a Teradata Extension pre V12 and an ANSI Version on Teradata V12.

### MERGE INTO Example that Matches

```
The first query shows Squiggy Jones in the Employee_Table with a salary of 32500.00.  
SELECT * FROM Employee_Table WHERE Employee_No = 2000000;  
Employee_No Dept_No Last_Name First_Name Salary  
-----  
2000000 I. Jones Squiggy 32500.00
```

I will now perform a MERGE that will have a MATCH

```
MERGE INTO Employee_Table  
USING VALUES (2000000, NULL, 'Jones', 'Squiggy', 40000.00)  
AS E (Emp, Dept, Las, Fir, Sal)  
ON Employee_No = Emp  
WHEN MATCHED THEN  
UPDATE set salary = sal  
WHEN NOT MATCHED THEN  
INSERT VALUES (Emp, Dept, Las, Fir, Sal);
```

```
SELECT * FROM Employee_Table WHERE Employee_No = 2000000;  
Employee_No Dept_No Last_Name First_Name Salary  
-----  
2000000 I. Jones Squiggy 40000.00
```

### MERGE INTO Example that does NOT Match

I will now perform a MERGE that will have a NOT MATCH situation because Employee\_No 3000000 does not exist in the Employee\_Table.

```
MERGE INTO Employee_Table  
USING VALUES (3000000, 400, 'Coffing', 'TeratTom', 300000.00)  
AS E (Emp, Dept, Las, Fir, Sal)  
ON Employee_No = Emp  
WHEN MATCHED THEN  
UPDATE set salary = sal  
WHEN NOT MATCHED THEN  
INSERT VALUES (Emp, Dept, Las, Fir, Sal);
```

```
SELECT * FROM Employee_Table WHERE Employee_No = 3000000;  
Employee_No Dept_No Last_Name First_Name Salary  
-----
```



**Teradata CoE****Teradata Lab Course**

```
3000000    400  coffing  Teratm  300000.00
```

There is no employee 3000000 that exists before the MERGE INTO statement, but once the MERGE INTO statement runs (and doesn't find a Match), it INSERTS employee 3000000 into the table.

**OReplace**

OReplace will replace characters or eliminate them.

```
Just give the replace string as a zero length string and it removes the character.
```

```
So if cont contains "123456789"
```

```
and you
```

```
OReplace (cont,"")
```

```
the result is
```

```
123456789
```

The OReplace function is a UDF that will replace certain characters with others.



## Date Functions



**Date Functions**



**Teradata CoE**      **Teradata Lab Course**

**Date, Time, and Current\_Timestamp Keywords**

```

SELECT
  DATE          AS "Date"
, CURRENT_DATE AS ANSI_Date
, TIME          AS "Time"
, CURRENT_TIME AS ANSI_Time
, CURRENT_TIMESTAMP(6) AS ANSI_Timestamp
  
```

Results:

```

Date      ANSI_Date    Time    ANSI_Time   ANSI_Timestamp
-----  -----  -----  -----  -----
08/13/2011 08/13/2011 17:48:25 17:48:25 08/13/2011 17:48:25.738000-04:00
  
```

Above are the keywords you can utilize to get the date, time, or timestamp. These are reserved words that the system will deliver to you when requested. The Keyword **TIMESTAMP** will fail. Notice the -04:00 at the end of the ANSI\_Timestamp. That is the time zone offset. This system is 4 hours behind Greenwich Mean Time (GMT).

Dates are stored internally as **INTEGERS** from a Formula

```

INTEGERDATE = ((Year *1000) + 10000) + (Month * 100) + Day
/* Example -Tom's Birthday January 10, 1959 */
INTEGERDATE = ((1959 * 1000) + 10000 + 100 * 100 + 10)
           + (Months * 100) + 590100
           + Day = 590110
           /* Year Portion
           * Month Portion
           * Day Portion */

/* Example -Tom's Birthday January 10, 1999 */
19990110
/* Example -Tom's Birthday January 10, 2000 */
1000110
/* Send Tom a birthday present on January 10, 2014 */
1140110
  
```

The reason the Smart Calendar works so well is that it stores **EVERY** date in Teradata as something known as an INTEGERTDATE.

Displaying Dates for INTEGERTDATE and ANSIDATE

```

SELECT Date          AS "Date"
, Current_Date     AS Display_Date
  
```

INTEGERTDATE (YYMMDD)      ANSIDATE (YYYY-MM-DD)

---

## Teradata CoE

## Teradata Lab Course

Date	Display Date	Date	Display Date
12/04/30	12/04/30	2011-04-30	2012-04-30

NEXUS Query Chameleon MM/DD/YYYY

Date Display Date

06/30/2012 06/30/2012

Teradata in release V2R3 defaulted to a display of YYMMDD. This is called the INTEGERDATE. This can be changed to ANSI DATE, which is YYYY-MM-DD for a specific session or by Default if the DBA changes the DATEFORM in DBS Control. This has nothing to do with how the date is stored internally. It has to do with the display of dates when using any ODBC tool or load utility. Above are some examples.

### DATEFORM

DATEFORM Controls the default display of dates.

DATEFORM display choices are either INTEGERDATE or ANSI DATE.

INTEGERDATE is (YYMMDD) and ANSI DATE is (YYYY-MM-DD)

DATEFORM is the expected format for import and export of dates in Load Utilities.

Can be over-ridden by USER or within a Session at any time.

The Default can be changed by the DBA by changing the DATEFORM in DBSControl.

INTEGERDATE (YYMMDD)	ANSI DATE (YYYY-MM-DD)
June 30, 2012	June 30, 2012

Teradata in release V2R3 defaulted to a display of YYMMDD. This is called the INTEGERDATE. This can be changed to ANSI DATE, which is YYYY-MM-DD for a specific session or by Default if the DBA changes the DATEFORM in DBS Control. This has nothing to do with how the date is stored internally. It has to do with the display of dates when using any ODBC tool or load utility.

Changing the DATEFORM in Client Utilities such as BTEQ



## Teradata CoE

## Teradata Lab Course

Enter your login or BTTEQ Command:  
Login local@dtc  
Password: \*\*\*\*\*  
Login successfully completed  
BTTEQ - Enter your DBC/SQL request or BTTEQ command:

SELECT DATE;

Notice the Word Date  
Date  
12/06/30 INTERGERDATE is the Default

BTTEQ - Enter your DBC/SQL request or BTTEQ command:  
SET Session DATEFORM = ANSI DATE; Changing the DATEFORM for this BTTEQ session.

SELECT DATE;  
Notice the Word Current Date  
Current Date  
2012-06-30 ANSIDATE is the Display Form

### Date, Time, and Timestamp Recap

```
SELECT Date AS "Date"  
, Current_Date AS ANSI_date  
INTERGERDATE (YYMMDD) ANSIDATE (YYYY-MM-DD)  
June 30, 2012 June 30, 2012
```

Date	ANSI_Date	Date	ANSI_Date
12/06/30	12/06/30	2012-06-30	2012-06-30

Dates are converted to an **integer** through a formula before being stored.

Dates are displayed by default as INTERGERDATE YYMM/DD.

The DBA can set up the system to display as **ANSIDATE YYYY-MM-DD**.

Keywords **Date** or **Current\_Date** will return the date automatically.

**Time**, **Current\_Time** and **Current\_Timestamp** are keywords to return time.

The **Nexus** Query Chameleon displays dates as **MMDDYYYY**.

### Timestamp Differences

```
SELECT Current_Timestamp(0) AS Col1  
, Current_Timestamp(6) AS Col2
```



**Teradata CoE**      **Teradata Lab Course**

---

**Answer Set**

Col1	Col2
2011/03/22 10:54:44	2011/03/22 10:54:44.123456
 Date	 Space
 Time	 Milliseconds

A timestamp has the date separated by a space and the time. In our second example, we have asked for 6 milliseconds.

**Finding the Number of Hours between Timestamps**

Load Test Data

```

Create Table Car_Maintenance
(
    Car_ID          INT,
    Current_Timestamp DATETIME,
    Interval       INTERVAL '7' DAY,
    Start_Timestamp DATETIME,
    End_Timestamp   DATETIME
)
;
CREATE INDEX IX_Car_Maintenance ON Car_Maintenance(Car_ID);
CREATE INDEX IX_Car_Maintenance_2 ON Car_Maintenance(Start_Timestamp);
CREATE INDEX IX_Car_Maintenance_3 ON Car_Maintenance(End_Timestamp);

INSERT INTO Car_Maintenance
SELECT Car_ID, Current_Timestamp, Interval '7' day, Current_Timestamp,
       Current_Timestamp - Interval '1' day, Current_Timestamp
FROM Car_Maintenance
WHERE Car_ID = 100;

```

```

SELECT Car_ID , Start_Timestamp, End_Timestamp
, (CAST(End_Timestamp AS DATE) -CAST(Start_Timestamp AS DATE)) * 24.0
+ (EXTRACT(HOUR FROM End_Timestamp) -EXTRACT(HOUR FROM Start_Timestamp)) * 1.0
+ (EXTRACT(MINUTE FROM End_Timestamp) -EXTRACT(MINUTE FROM Start_Timestamp)) / 60.0
+ (EXTRACT(SECOND FROM End_Timestamp) -EXTRACT(SECOND FROM Start_Timestamp)) / 3600.0
AS Hours
FROM Car_Maintenance ORDER BY Car_ID ;

```

Car_ID	Start_Timestamp	End_Timestamp	Hours
1	08/07/2013 9:14:19.000000	08/14/2013 9:14:19.000000	168.000000
2	08/13/2013 9:11:11.640000	08/14/2013 9:11:11.640000	24.000000
3	08/11/2013 9:18:12.370000	08/14/2013 9:18:12.370000	72.000000
4	08/11/2013 9:18:39.940000	08/14/2013 9:18:39.940000	72.000000

The above example is how you find the number of hours between Timestamps.

**Troubleshooting Timestamp**

```

SELECT Timestamp(0) AS Col0
, Timestamp(6) AS Col0

```

Error

There is Date and Current\_Date (both work).

There is Time and Current\_Time (both work).

There is **NO** Timestamp, but only **Current\_Timestamp**.

There is NO Timestamp KEYWORD, but only ANSI's Current\_Timestamp!

**Add or Subtract Days from a date**

```

SELECT Order_Date
,Order_Date + 60 as "Due Date"
,Order_Total

```

---

 Capgemini  
CONSULTING. INNOVATION. EXECUTING.

```

    "Due date" -10 as Discount
    FROM ...
    Order_Table
    ORDER BY i ;

```

Order Date	Due Date	Order Total	Discount	Discounted
05/04/1999	07/03/1999	12347.53	04/21/1999	12100.57
01/01/1999	03/02/1999	8005.91	02/20/1999	7845.79
10/01/1999	11/30/1999	12345.67	10/20/1999	12345.74
10/01/1999	11/30/1999	5111.47	11/20/1999	5009.24
10/10/1999	12/09/1999	18231.62	11/29/1999	18226.99

When you add or subtract from a Date you are adding/subtracting Days

Because Dates are stored internally on disk as integers, it makes it easy to add days to the calendar. In the query above, we are adding 60 days to the Order\_Due

#### A Summary of Math Operations on Dates

- ➊ DATE - DATE = Interval (days between dates)
- ➋ DATE + or - Integer = Date

Let's find the number of days Tera-Tom has been alive since his 2014 birthday.

```

SELECT (1140110(date)) - (590110 (date)) (Title 'Tera-Tom's Age In Days') ;
Tera-Tom's Age In Days
2009

```

Below is the same exact query, but with a clearer example of the dates.

```

SELECT ("2014-01-10" (date)) -("1959-01-10" (date)) (Title 'Tera-Tom's Age In Days') ;
Tera-Tom's Age In Days
2009

```

A DATE - DATE is an interval of days between dates. A DATE + or - Integer = Date. Both queries above perform the same function; but the top query uses the internal date functions, and the query on the bottom does dates the traditional way.

#### Using a Math Operation to find your Age in Years

- ➊ DATE - DATE = Interval (days between dates)
- ➋ DATE + or - Integer = Date

Let's find the number of days Tera-Tom has been alive since his 2014 birthday.

```

SELECT (1140110(date)) - (590110 (date)) (Title 'Tera-Tom's Age In Days') ;
Tera-Tom's Age In Days
2009

```

Let's find the number of years Tera-Tom has been alive since his 2014 birthday

```

SELECT ((1140110(date)) - (590110 (date))) / 365 (Title 'Tera-Tom's Age In Years') ;
Tera-Tom's Age In Days
55

```



## Teradata CoE

## Teradata Lab Course

A DATE - DATE is an interval of days between dates. A DATE + or - Integer = Date. Both queries above perform a Date function; but the top query brings back Tom's age in days, and the bottom query brings back Tom's age in years.

Find What Day of the week you were Born

Let's find the actual day of the week Tom-Tom was born

```
SELECT 'Tom-Tom was born on day ' || ((59010(date))-(101(date))) MOD 7 (TITLE '),
       Tom-Tom was born on day 5
      FROM DUAL;
```

Result	Day of the Week
0	Monday
1	Tuesday
2	Wednesday
3	Thursday
4	Friday
5	Saturday
6	Sunday

The above subtraction results in the number of days between the two dates. Then the MOD 7 divides by 7 to get rid of the number of weeks and results in the remainder. A MOD 7 can only result in values 0 thru 6 (answers 1 less than the MOD operator). Since January 1, 1900 (101(date)) is a Monday, Tom was born on a Saturday.

The ADD\_MONTHS Command

```
Order_Table
Order_Number Customer_Number Order_Date Order_Total
123456           11111111 1998/05/04    12347.50
123512           11111111 1998/01/01    8005.92
123523           11111111 1998/05/15    12347.50
123585           97323456 1998/10/10   15231.42
123771           57998932 1998/09/08   20854.00

SELECT Order_Date
      ,ADD_MONTHS(Order_Date,2) as "Due Date"
      ,Order_Total
  FROM Order_Table ORDER BY 1;
```

This is the Add\_Months Command. What you can do with it is add a month or many months to your column date. Can you convert this to one year?

Using the ADD\_MONTHS Command to Add 1 Year

```
Order_Table
Order_Number Customer_Number Order_Date Order_Total
123456           11111111 1998/05/04    12347.50
123512           11111111 1998/01/01    8005.92
```



## Teradata CoE

## Teradata Lab Course

```
123852      31323134  1999/10/01    5111.47  
123853      31323134  1999/10/01    12347.53  
123777      37894683   1999/09/09    23454.84
```

```
SELECT Order_Date  
      ,Add_Months(Order_Date,12) as "Due Date"  
      ,Order_Total  
  FROM  Order_Table  
 ORDER BY 1;
```

There is no Add\_Year command, so put in 12 months for 1-year

The Add\_Months command adds months to any date. Above we used a great technique that would give us 1 year. Can you give me 5 years?

Using the ADD\_MONTHS Command to Add 5 Years

Order_Number	Customer_Number	Order_Date	Order_Total
123854	11111111	1999/05/04	12347.53
123855	11111111	1999/05/04	12347.53
123852	31323134	1999/10/01	5111.47
123778	31323134	1999/10/01	12347.53
123777	37894683	1999/09/09	23454.84

```
SELECT Order_Date  
      ,Add_Months(Order_Date,12 * 5) as "Due Date"  
      ,Order_Total  
  FROM  Order_Table  
 ORDER BY 1;
```

In this example we multiplied 12 months times 5 for a total of 5 years!

Above you see a great technique for adding multiple years to a date. Can you now SELECT only the orders in September?

The EXTRACT Command

Order_Number	Customer_Number	Order_Date	Order_Total
123854	11111111	1999/05/04	12347.53
123855	11111111	1999/05/04	12347.53
123852	31323134	1999/10/01	5111.47
123778	31323134	1999/10/01	12347.53
123777	37894683	1999/09/09	23454.84

```
SELECT Order_Date  
      ,Add_Months(Order_Date,12 * 5) as "Due Date"  
      ,Order_Total  
  FROM  Order_Table  
 WHERE  EXTRACT(Month from Order_Date) = 9  
 ORDER BY 1;
```

The EXTRACT command extracts portions of Date, Time, and Timestamp.



This is the Extract command. It extracts a portion of the date. It can be used in the SELECT list, the WHERE Clause, or the ORDER BY Clause!

#### EXTRACT from DATES and TIME

```
SELECT Current_Date
      ,EXTRACT(YEAR from Current_Date) as Yr
      ,EXTRACT(MONTH from Current_Date) as Mo
      ,EXTRACT(DAY from Current_Date) as Da
      ,Current_Time
      ,EXTRACT(HOUR from Current_Time) as Hr
      ,EXTRACT(MINUTE from Current_Time) as Mn
      ,EXTRACT(SECOND from Current_Time) as Ss
      ,EXTRACT(MILLISECOND from Current_Time) as Tb
      ,EXTRACT(TIMEZONE_MINUTE from Current_Time) as Tm;
```

Answer Set

Current_Date	Yr	Mo	Da	Current_Time	Hr	Mn	Ss	Tb	Tm
2013/03/22	2013	03	22	2013-03-22 20:01:42.000000	20	01	42	000000	0

Just like the Add\_Months, the EXTRACT Command is a Temporal Function or a Time-Based Function.

#### CURRENT\_DATE and EXTRACT or CURRENT\_DATE and Math

```
SELECT Current_Date
      ,EXTRACT(YEAR from Current_Date) as Yr
      ,EXTRACT(MONTH from Current_Date) as Mo
      ,EXTRACT(DAY from Current_Date) as Da
      ,Current_Date +100 as CurrDate
      ,(Current_Date +100)/100 as ModMath
      ,Current_Date Mod 100 as DayMod;
```

Math can be used to extract portions of a Date!

Answer Set

Current_Date	Yr	Mo	Day	ModMath	DayMod	
2013/03/22	2013	03	22	2013	03	22

The Extract Temporal Function can be used to extract a portion of a date. As you can see, Basic Arithmetic accomplishes the same thing.

#### CAST the Date of January 1, 2011 and the Year 1800

```
SELECT
      cast('2011-01-01' as date) as ADT_Literal
      ,cast('2011-01-01' as date) as INTDT_literal
      ,cast('11-01-01' as date) as YY_Literal
      ,cast('11-01-01' as date) as Date_Stored
      ,cast(Date '1800-01-01' as Integer) as IntDate_1800;
```

Answer Set

ADT_Literal	INTDT_literal	YY_Literal	Date_Stored	IntDate_1800
2011-01-01	2011-01-01	01/01/1800	01/01/1800	-999999

The Convert and Store (CAST) command is used to give columns a different data type temporarily for the life of the query.



## Teradata CoE

## Teradata Lab Course

Notice our dates and how they're stored.

### The System Calendar

Teradata systems have a [table](#) called [Caldates](#).

[Caldates](#) has only one column in it called [Caldates](#).

[Caldates](#) is a date column that contains a row for each date starting from January 1, 1900 to December 31, 2100.

No user can access the table [Caldates](#) directly.

Views in the [Sys\\_Calendar](#) database accesses [Caldates](#).

A [view](#) called [Calendar](#) is how USERS work with the calendar.

Users use [Sys\\_Calendar.Calendar](#) for advanced dates.

In every Teradata system, there is something known as a System Calendar (or as Teradata calls it "Sys\_Calendar.Calendar"). Get ready for AWESOME!

### Using the System Calendar in its Simplest Form

```
SELECT * FROM Sys_Calendar.Calendar WHERE Calendar_Date = '1959-01-17';
```

```
Calendar_Date = '01/10/1959'  
day_of_week = 7 (Sunday = 1)  
day_of_month = 10  
day_of_year = 10  
is_leap_year = 1  
weekday_of_month = 2 (since Jan 1, 1900)  
week_of_month = 1 (for partial week for any month not starting with Sunday)  
week_of_calendar = 3079 (since Jan 1, 1900)  
month_of_year = 1  
month_of_calendar = 709 (since Jan 1, 1900)  
quarter_of_year = 1  
quarter_of_calendar = 237 (since Jan 1, 1900)  
century_of_calendar = 1959
```

Tera-Tom was born on a Saturday! It was the first full week of the month, the first full week of the year, and it was the first quarter of the year!

How to really use the [Sys\\_Calendar.Calendar](#)



```
DATE '1999-01-10' is stored as 990110  
DATE '2000-01-10' is stored as 1000110
```

4 bytes store Date\_col internally because dates are considered a 4-byte integer.

#### Storing Time Internally

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,  
NO BEFORE JOURNAL,  
NO AFTER JOURNAL,  
CHECKSUM = DEFAULT  
(Date_col  
TIME_col  
TIMETIMEZONE_col TIME(6) WITH TIME ZONE,  
TIMESTAMP_col TIMESTAMP(6),  
TIMEZONE_col TIMESTAMP(6) WITH TIME ZONE)  
UNIQUE PRIMARY INDEX (TIMEZONE_col);
```

```
Time(n) stored as HHMMSS.nnnnnn
```

It takes 6 bytes to store Time\_col internally.

#### Storing TIME with TIME ZONE Internally



**Teradata CoE****Teradata Lab Course**

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,
  NO BEFORE JOURNAL,
  NO AFTER JOURNAL,
  CHECKSUM = DEFAULT
  (Date_col          Date,
   TIME_col           TIME(6),
   TIMETIMEZONE_col  TIME(6) WITH TIME ZONE,
   TIMESTAMP_col     TIMESTAMP(6),
   TIMEZONE_col      TIMESTAMP(6) WITH TIME ZONE)
  UNIQUE PRIMARY INDEX (TIMEZONE_col);
```

Time(n) WITH ZONE stored as HHMMSS.nnnnnn+HHM

It takes 8 bytes to store TimeTimezone\_col internally.

Storing Timestamp Internally

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,
  NO BEFORE JOURNAL,
  NO AFTER JOURNAL,
  CHECKSUM = DEFAULT
  (Date_col          Date,
   TIME_col           TIME(6),
   TIMETIMEZONE_col  TIME(6) WITH TIME ZONE,
   TIMESTAMP_col     TIMESTAMP(6),
   TIMEZONE_col      TIMESTAMP(6) WITH TIME ZONE)
  UNIQUE PRIMARY INDEX (TIMEZONE_col);
```

Timestamp(n) stored as YYMMDDHHMMSS.nnnnnn

It takes 10 bytes to store TimeStamp\_col internally.

Storing Timestamp with TIME ZONE Internally



```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,
NO BEFORE JOURNAL,
NO AFTER JOURNAL,
CHECKSUM = DEFAULT
(Date_col          Date,
TIME_col           TIME(6),
TIMEZONE_col      TIME(6) WITH TIME ZONE,
TIMESTAMP_col     TIMESTAMP(6),
TIMEZONE_COL      TIMESTAMP(6) WITH TIME ZONE)
UNIQUE PRIMARY INDEX (TIMEZONE_col);
```

TimeStamp(n) With Zone stored as YYMMDDHHMMSS.nnnnnn+HHMM

It will take 12 bytes to store Timezone\_col internally.

#### Storing Date, Time, and Timestamp with Zone Internally

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,
NO BEFORE JOURNAL,
NO AFTER JOURNAL ,
CHECKSUM = DEFAULT
(Date_col          Date,
TIME_col           TIME(6),
TIMEZONE_COL      TIME(6) WITH TIME ZONE,
TIMESTAMP_col     TIMESTAMP(6),
TIMEZONE_COL      TIMESTAMP(6) WITH TIME ZONE)
UNIQUE PRIMARY INDEX (TIMEZONE_col);
```



Date	Stored Internally	4 Bytes
Time(n)	Stored Internally	8 Bytes
Timestamp	Stored Internally	10 Bytes
Timestamp with zone	Stored Internally	12 Bytes

Each data type increases its internal storage by 2 bytes.

#### Time Zones

A time zone relative to London(UTC) might be:
DA-----+---Miami-----+Frankfurt-----+Hong Kong
+03:00            -04:00            -02:00
A time zone relative to New York (EST) might be:
DA-----+---Miami-----+Frankfurt-----+Hong Kong
+01:00            -05:00            -13:00

Time zones are set either at the **system level** (DBS Control), the **user level** (when user is created or modified), or at the **session level** as an override.

Teradata has the ability to access and store both the hours and the minutes reflecting the difference between the world time zone and the system time zone. From a World perspective, the difference is normally the number of hours between a specific location in Latin America and the United Kingdom location that was historically called Greenwich Mean Time (GMT). Since the Greenwich observatory has been decommissioned, the new reference to this same time zone is called Universal Time Coordinate (UTC).

#### Setting Time Zones

A Time Zone should be established for the system and every user in each different time zone.

Setting the **system default** time zone is done by the DBA in the DBSControl record:

HOSTTIMEGENERAL 17 = # /\* Hours,      # = +1 to -15 \*/

HOSTTIMEINTERVAL 17 = # /\* Minutes,      # = +1 to -59 \*/

Setting a **User's** time zone requires choosing either **LOCAL**, **NULL**, or an **explicit value**:

```
CREATE USER TexaTime
TIME ZONE = LOCAL /* use system level */
              /* or choose to use session level at logon */
              /* or explicit setting */
              /* +16:00 */
              /* -16:00 */
              /* -04:30 */
```

Setting a **Session's** time zone:

```
SET TIME ZONE LOCAL /* use system level */
SET TIME ZONE NULL /* use session level */
SET TIME ZONE INTERVAL '08:00' HOUR TO MINUTE /* explicit setting */
```

A Teradata session can modify the time zone without requiring a logoff and logon.

#### Seeing your Time Zone



**Teradata CoE**      **Teradata Lab Course**

---

Help Session:

User	Account	Logon	Logon	Current	Collation	Char	Session	Transaction	Current	Differences	Session
Name	Name	Date	Time	Database	Set	Schemas	Set	Isolation	Deferred	Time Zone	
DbC	DbC	12/06/17	15:55:39	SQ_CCLASS	ASCI	ASCI	Default	ReadCommitted	00:00	IntegerDate	

Not all input  
is displayed  
above from the  
HELP.Session

A user's time zone is now part of the information maintained by Teradata. The settings can be seen in the extended information available in the HELP SESSION request. Teradata converts all TIME and TIMESTAMP values to Universal Time (UTC) prior to storing them. All operations, including hashing, collation, and comparisons that act on TIME and TIMESTAMP values, are performed using their UTC forms. This will allow users to CAST the information to their local times.

**Creating a Sample Table for Time Zone Examples**

```
CREATE TABLE Tstamp_Test
[ TS_Locn CHAR(10)
, TS_Tzname TIMESTAMP(6) WITH TIME ZONE
, TS_Without_Zone TIMESTAMP(6)
]
UNIQUE PRIMARY INDEX ( TS_Zone );
```

A user's time zone is now part of the information maintained by Teradata. The settings can be seen in the extended information available in the HELP SESSION request.

**Inserting Rows in the Sample Table for Time Zone Examples**

```
Enter your logon or BTEQ Command:
Logon locn 3514773898
Logon successfully completed
BTEQ - Enter your DB2/SQl request or BTEQ command:
INSERT INTO Tstamp_Test ('EST', timestamp '2000-10-01 08:12:00',
                           timestamp '2000-10-01 08:12:00')
SET TIME ZONE INTERVAL '03:00' HOUR TO MINUTE ;
  INSERT INTO Tstamp_Test ('UTC', timestamp '2000-10-01 08:12:00',
                           timestamp '2000-10-01 08:12:00')
SET TIME ZONE INTERVAL '-03:00' HOUR TO MINUTE ;
  INSERT INTO Tstamp_Test ('PST', timestamp '2000-10-01 08:12:00',
                           timestamp '2000-10-01 08:12:00')
SET TIME ZONE INTERVAL '-11:00' HOUR TO MINUTE ;
  INSERT INTO Tstamp_Test ('HKT', timestamp '2000-10-01 08:12:00',
                           timestamp '2000-10-01 08:12:00')
```

**Selecting the Data from our Time Zone Table**

```
SELECT * FROM Tstamp_Test ;
```

---

 Capgemini  
PERFORMING BUSINESS BETTER

## Teradata CoE

## Teradata Lab Course

TS_Zone	TS_With_Zone	TS_Without_Zone
UTC	2000-10-01 08:12:00.000000+00:00	2000-10-01 08:12:00.000000
EST	2000-10-01 08:12:00.000000-04:00	2000-10-01 08:12:00.000000
PST	2000-10-01 08:12:00.000000-07:00	2000-10-01 08:12:00.000000
HKT	2000-10-01 08:12:00.000000+11:00	2000-10-01 08:12:00.000000

Notice the Accompanying Time Zone Offsets

Our Insert statements were done at 08:12:00 exactly. Notice the Time Zone offsets in the column TS\_With\_Zone and how they're not there for the column TS\_Without\_Zone. Teradata converts all TIME and TIMESTAMP values to Universal Time Coordinated (UTC) prior to storing them. All operations, including healing, collation, and comparisons that act on TIME and TIMESTAMP values, are performed using their UTC forms. This will allow users to CAST the information to their local times.

Normalizing our Time Zone Table with a CAST

```
SELECT TS_Zone, TS_With_Zone
,CAST(TS_With_Zone AS TIMESTAMP(8)) AS T_Norma
FROM Tstamp_Text ORDER BY 3 ;
```

TS_Zone	TS_With_Zone	T_Norma
UTC	2000-10-01 08:12:00.000000+00:00	2000-10-01 08:12:00.000000
EST	2000-10-01 08:12:00.000000-04:00	2000-10-01 08:12:00.000000
PST	2000-10-01 08:12:00.000000-07:00	2000-10-01 08:12:00.000000
HKT	2000-10-01 08:12:00.000000+11:00	2000-10-01 08:12:00.000000

The System is on EST Time. The New Times are Normalized to the time zone of the System!

Notice that the Time Zone value was added to, or subtracted from, the time portion of the timestamp to adjust it to a perspective of the same time zone. As a result, it has normalized the different Time Zones in respect to the system time.

As an illustration, when the transaction occurred at 08:12 AM locally in the PST Time Zone, it was already 11:12 AM in EST, the location of the system. The times in the columns have been normalized in respect to the time zone of the system.

Intervals for Date, Time and Timestamp

Interval Chart	
Simple Intervals	More involved Intervals
YEAR	DAY TO HOUR
MONTH	DAY TO MINUTE
DAY	DAY TO SECOND
HOUR	HOUR TO MINUTE
MINUTE	HOUR TO SECOND
SECOND	MINUTE TO SECOND

To make Teradata SQL more ANSI compliant and compatible with other RDBMS SQL, Teradata has added INTERVAL processing. Intervals are used to perform DATE, TIME and TIMESTAMP arithmetic and conversion.

Although Teradata allowed arithmetic on DATE and TIME, it was not performed in accordance to ANSI standards and



therefore, an extension instead of a standard. With INTERVAL being a standard instead of an extension, more SQL can be ported directly from an ANSI compliant database to Teradata without conversion.

#### Interval Data Types and the Bytes to Store Them

Interval Chart

Bytes	Data Type	Comments
2	INTERVAL YEAR	
4	INTERVAL YEAR TO MONTH	
2	INTERVAL MONTH	
2	INTERVAL DAY TO HOUR	
2	INTERVAL DAY	
8	INTERVAL DAY TO MINUTE	
10/12	INTERVAL DAY TO SECOND	10 for 32-bit systems; 12 for 64-bit
4	INTERVAL HOUR TO MINUTE/4	
4	INTERVAL HOUR TO SECOND/4	
2	INTERVAL MINUTE/2	
6/8	INTERVAL MINUTE TO SECOND	6 for 32-bit systems; 8 for 64-bit
1	INTERVAL SECOND	6 for 32-bit systems; 8 for 64-bit

#### The Basics of a Simple Interval

```
SELECT Current_Date as Our_Date,
       Current_Date + Interval '1' Day as Plus_1_Day
      ,Current_Date + Interval '3' Month as Plus_3_Months
      ,Current_Date + Interval '5' Year as Plus_5_Years;
```

In the example SQL above, we take a single date and add 1 day, 3 months, and 5 years. Notice that our current\_date is 06/19/2012 and that our intervals come out precisely.

#### Troubleshooting the Basics of a Simple Interval

```
SELECT Date '2012-01-29' as Our_Date
      ,Date '2012-01-29' + INTERVAL '1' Month as Leap_Year;
```

Out Date      Leap\_Year

01/29/2012    02/28/2012

```
SELECT Date '2011-01-29' as Our_Date
      ,Date '2011-01-29' + INTERVAL '1' Month as Leap_Year;
```

Error-Invalid Date

The first example works because we added 1 month to the date 2012-01-29 and we got 2012-02-29. Because this was a leap year, there actually is a date of February 29, 2012. The next example is the real point. We have a date of 2011-01-29 and we add 1 month to that. But there is no February 29<sup>th</sup> in 2011, so the query fails.

#### Interval Arithmetic Results

DATE and TIME arithmetic Results using intervals:



```

DATE - DATE      = Interval
TIME - TIME      = Interval
TIMESTAMP - TIMESTAMP = Interval
DATE - :c1 + Interval = DATE
TIME - :c1 + Interval = TIME
TIMESTAMP - :c1 + Interval = TIMESTAMP
Interval - :c1 + Interval = Interval

```

To use DATE and TIME arithmetic, it is important to keep in mind the results of various operations. The above chart is your Interval guide.

#### A Date Interval Example

```
SELECT (DATE '1999-10-01' - DATE '1988-10-01') DAY AS Actual_Days;
```

**ERROR - Interval Field Overflow**

The error occurred because the default  
for all intervals is 2 digits.

```
SELECT (DATE '1999-10-01' - DATE '1988-10-01')DAY(4) AS Actual_Days;
```

Makes the output 4 digits

Actual_Days
4017

The default for all intervals is 2 digits. We received an overflow error because the Actual\_Days is 4017. The second example works because we declared the output to be 4 digits (the maximum for intervals).

#### A Time Interval Example

```

SELECT ((TIME '12:45:01' + TIME '10:10:01') HOUR AS Actual_Hours
       ,(TIME '12:45:01' - TIME '10:10:01')MINUTE(0) AS Actual_Minutes
       ,(TIME '12:45:01' - TIME '10:10:01')SECOND(0) AS Actual_Seconds
       ,(TIME '12:45:01' - TIME '10:10:01')SECOND(4) AS Actual_Seconds1

```

Actual_Hours	Actual_Minutes	Actual_Seconds	Actual_Seconds1
2	155	9300.000000	9300.0000

```

SELECT ((TIME '12:45:01' - TIME '10:10:01')HOUR AS Actual_Hours
       ,(TIME '12:45:01' - TIME '10:10:01')MINUTE AS Actual_Minutes
       ,(TIME '12:45:01' - TIME '10:10:01')SECOND(4) AS Actual_Seconds
       ,(TIME '12:45:01' - TIME '10:10:01')SECOND(4.4) AS Actual_Seconds1

```

**ERROR - Interval Field Overflow**



The default for all intervals is 2 digits, but notice in the top example we put in 3 digits for Minute, 4 digits for Second, and 4,4 digits for the Actual\_Seconds4. If we had not, we would have received an overflow error as in the bottom example.

#### A - DATE Interval Example

```
SELECT Current_Date,
       INTERVAL '+2' YEAR + CURRENT_DATE as Two_years_Ago;
Date          Two_Years_Ago
06/19/2012   06/19/2010
```

The above Interval example uses a “Z” to go back in time.

#### A Complex Time Interval Example using CAST

Below is the syntax for using the CAST with a date:  
`SELECT CAST (interval_val AS INTERVAL <interval> )  
FROM table-name; /`

The following converts an INTERVAL of 6 years and 2 months to an INTERVAL number of months:

```
SELECT CAST( (INTERVAL '6-02' YEAR TO MONTH) AS INTERVAL MONTH );
          6-02
          74
```

The CAST function (Convert and Store) is the ANSI method for converting data from one type to another. It can also be used to convert one INTERVAL to another INTERVAL representation. Although the CAST is normally used in the SELECT list, it works in the WHERE clause for comparison purposes.

#### A Complex Time Interval Example using CAST

This request attempts to convert 1300 months to show the number of years and months. Why does the first example fail, but the second find success?

```
SELECT CAST(INTERVAL '1300' MONTH AS INTERVAL YEAR TO MONTH) as YMo;
          YMo
          ERROR

SELECT CAST(INTERVAL '1300' MONTH AS interval YEAR(3) TO MONTH) as YMo;
          YMo
          108-04
```

The top query failed because the INTERVAL result defaults to 2-digits and we have a 3-digit answer for the year portion (108). The bottom query fixes that by specifying 3-digits. The biggest advantage in using the INTERVAL processing is that SQL written on another system is now compatible with Teradata.

#### The OVERLAPS Command

Compatibility: Teradata Extension

The syntax of the OVERLAPS command is:

```
SELECT <literal>
      WHERE <(start-date-time), <(end-date-time)> OVERLAPS
            <(start-date-time), <(end-date-time)> /
```



```
SELECT 'The Dates Overlap' (TITLE )
WHERE (DATE '2001-01-01', DATE '2001-11-30')
OVERLAPS (DATE '2001-10-15', DATE '2001-12-31')
```

Answer → The Dates Overlap

When working with dates and times, sometimes it is necessary to determine whether two different ranges have common points in time. Teradata provides a Boolean function to make this test for you. It is called OVERLAPS; it evaluates true if multiple points are common, otherwise it returns a false. The literal is returned because both date ranges have from October 15 through November 30 in common.

An OVERLAPS Example that Returns No Rows

```
SELECT 'The dates overlap' (TITLE )
WHERE (DATE '2001-01-01', DATE '2001-11-30')
OVERLAPS (DATE '2001-11-30', DATE '2001-12-31')
```

Answer → No rows found

The above SELECT example tests two literal dates and uses the OVERLAPS to determine whether or not to display the character literal.

The literal was not selected because the ranges do not overlap. This means the common single date of November 30 does not contain an overlap. When dates are used, 2 days must be involved. When time is used, 2 seconds must be contained in both ranges.

The OVERLAPS Command using TIME

```
SELECT 'The Times Overlap' (TITLE )
WHERE (TIME '03:00:00', TIME '02:00:00')
OVERLAPS (TIME '02:01:00', TIME '04:15:00')
```

Answer → The Times Overlap

The above SELECT example tests two literal time and uses the OVERLAPS to determine whether or not to display the character literal.

This is a good example, and it is shown to prove a point. At first glance, it appears as if this answer is incorrect because 02:01:00 looks like it starts 1 second after the first range ends. However, the system works on a 24-hour clock when a date and time (timestamp) are not used together. Therefore, the system considers the earlier time of 2AM as the start, and the later time of 5 AM as the end of the range. Therefore, not only do they overlap, the second range is entirely contained in the first range.

The OVERLAPS Command using a NULL Value

```
SELECT 'The Times Overlap' (TITLE )
WHERE (TIME '10:00:00', NULL)
```



```
OVERLAPS (TIME '01:01:00', TIME '04:15:00'))
```

Answer

→ No Rows Found

The above SELECT example tests two literal dates and uses the OVERLAPS to determine whether or not to display the character literal.

When using the OVERLAPS function, there are a couple of situations to keep in mind:

1. A single point in time, i.e. the same date, does not constitute an overlap. There must be at least one second of time in common for TIME, or one day when using DATE.
2. Using a NULL as one of the parameters, the other DATE or TIME constitutes a single point in time instead of a range.

